

What I found when I stopped agreeing with us. Twenty-six research rounds, 109 reports, and a twelve-part adversarial audit against our own case. The editor is the product — that decision is yours and it is made. This is about what the editor should *do*, what it should say, and what it should cost. Anything unmarked I checked against a primary source myself. Where our own record was wrong, I have said so.

KEY FIGURES

4 / 12,556

complaints about mangled formatting, the thing we planned to sell

89

context-pack products launched on HN in 20 months. Median 2 points

₹1.09L

monthly contribution from the 2-year success case. Same as 78 hours of consulting

23.7%

of all AI complaints are slop. Growing +149%. Nobody is on it

62%

of our own load-bearing claims were wrong when someone opened a source

PART I

Start here

1 The one-paragraph version

- **The engine is right. Everything we said about it was wrong.**
- Byte-exactness is real, rare and correct — and users do not talk about it. Four complaints in 12,556.
- Our flagship demo, per-hunk accept/reject, is already shipped by Zed. Verbatim, in their docs.
- Our wedge, generating handovers and context packs, has been launched 89 times in 20 months. Median score: 2 points.
- **What survives is the mechanism, pointed at a different problem.** The engine records byte ranges. Slop is 23.7% of complaints and growing 149% a year. Nobody can tell which parts of a document a machine wrote. We can.
- **The product is still a markdown editor.** The USP is not that it preserves your bytes. It is that it is the only editor that knows which bytes are yours.

RISK

The uncomfortable arithmetic, and I am not going to bury it. The two-year success case — 502 paying users, 171,200 cumulative visitors — produces ₹1,09,135/month. That is the same money as 78 hours of consulting at your own measured rate, billable starting Monday. Building the product is not the fastest path to that number. It is the path to owning an asset. Decide with that in front of you, not behind you.

2 What changed, in one table

EXHIBIT 1

Every belief we tested, and what the evidence returned

#	What we believed	Verdict	What the source actually says
1	Byte-exactness makes people switch	REFUTED	4 of 12,556 HN comments · 1 of 3,749 launch comments · once in 1,000 Reddit posts (0.1%) · Line endings in 1 of 43,656 GitHub issues · one switch in 11 years
2	Per-hunk review is our demo	REFUTED	Zed docs: "accept or reject each individual change hunk, or the whole set" — I opened the page myself
3	Handovers and context packs are the wedge	REFUTED	89 launches in 20 months, median 2 points, 88 of 89 never reached 50. Four were called <i>Handoff</i>
4	Decision-flow renders are the obvious first render	REFUTED	563 of 143,283,562 Obsidian downloads. 0.0004%. 3 plugins of 7,079. 0 of 1,600 Reddit posts
5	₹299 is a fair price	REFUTED	Obsidian gives the editor away and charges \$4 for sync alone. We priced a whole editor at \$3.13
6	The engine is the differentiator	CONFIRMED	Cursor staff, three days before I looked: destroyed carriage returns on CRLF is "a known issue we're tracking"
7	The correctness work is real	CONFIRMED	8,513-file pinned corpus, byte-identical verification, 1,575 tests
8	Sync is out of scope	REFUTED	#1 loved feature, #3 hated, #1 switching trigger, 1,251 of 43,656 issues

GOOD

Read row 6 and 7 together. The incumbents have the defect we fixed, they know it, and they have not fixed it. That is a real edge. We were just selling it with the wrong sentence.

PART II

The critique

3 Product and premise

The deepest finding, and it is structural.

- **We have two mutually exclusive buyers and have never picked one.** Every plan in the record serves the AI-native developer. Every problem statement that survives our own filtering serves a liable non-developer — someone who gets in trouble when a document is wrong.
- Those two people want different products. One wants speed. The other wants proof.
- The record reads as one product because nobody made us choose. **Choose.**

On the chain of reasoning.

Link	Holds?	Why
Markdown is the right substrate	YES	Files, diffs, git, and every model is trained on it
Byte-exact editing is the differentiator	PARTIAL	True and unfelt. It is a mechanism, not a message
People building software with AI are the buyer	OPEN	They are the loudest. They may not be the ones who pay
An editor is the right form factor	YES	Your decision, and defensible — it is where the work happens
A subscription is the right model	NO	The category gives the editor away and charges for services

- **Sunk cost is the biggest bias in the room.** 2.2 million words of documentation against 25,407 lines of code — 87 words per line. That ratio is what someone does when they are uncertain and writing feels like progress.
- **"Refuse rather than guess" is an engineering aesthetic that happens to also be right.** Do not confuse the two. Users do not buy refusal. They buy the confidence it produces.

4 Features and usability

EXHIBIT 2

The features, graded honestly

Feature	Who uses it, how often	Verdict
Per-hunk accept/reject on AI edits	Every AI user, daily	SHIP — but Zed has it. Not a differentiator, a table stake
Provenance: which bytes the machine wrote	Everyone reviewing AI output	SHIP — nobody has this. This is the product
Vault-wide refactor with reviewable diff	Anyone with >100 notes, monthly	SHIP — 86 likes on broken-links-on-rename
Nested-construct live preview	Every Obsidian user, constantly	LATER — 501 likes, the most-voted bug in the category. Free wedge
Sync	Everyone, daily	LATER — #1 loved, and the only proven price in the category
Degradation certificate, live	Almost nobody	KILL — technically lovely, commercially irrelevant
Decision-flow renders	0.0004% of the market	KILL — the least-wanted render measurable
Byte-level time travel	Nobody asked	KILL
Two-way renders	Nobody asked	KILL

Usability problems nobody has solved, including us.

- **Diff review is work.** We treat it as a feature. A 40-hunk change is a chore, not a delight. If our review is not faster than `git add -p` — which is free — we have added a step.
- **Refusal is a "no" to a paying user.** Model the moment: the editor says it cannot address something unambiguously. What does the user do? If the answer is "edit it by hand", we removed nothing.

- **The table-stakes gap is real.** A reviewer who cannot find quick-switch stops writing the review. That gap ranks above every novel feature on this list.

5 Business and reachability

Pricing is wrong on the axis, not the number.

- Obsidian: **editor free, sync \$4/user/month, publish per site.** I opened the page.
- We planned to charge for the editor. In this category that is the one thing nobody pays for.
- ₹299 is not "cheap". It is **below the credibility floor** — a whole editor for less than one competitor's single add-on reads as unserious.

EXHIBIT 3

The three routes to revenue, scored

Route	Evidence of demand	Time to revenue	Reuses the engine	Verdict
Services, productised	9/10	This month	~10%	Highest expected value. Not an asset
Provenance in the editor	9/10	~5 months	~90%	The build
Sync	9/10	~12 months	~60%	The only proven price. Hard, later

Reachability — and this is our weakest area.

- Every channel we have identified is **borrowed**: HN is one shot, the Obsidian plugin lives under someone else's policy, and our brand term is the substrate's generic name.
- **We own nothing.** No list, no audience, no distribution.
- The record has never answered *"what do you do on Monday morning to get the first hundred users?"* — and a channel list is not an answer.

WARN

Virality does not come free with a good product. The things that spread in this category are: a free tool people recommend, a visible artefact with a mark on it, and a plugin in someone else's store. We currently plan none of them.

6 Execution and capacity

- **The record contradicts itself on the most important number.** This plan says two founders. §53 says *"One person. This is the binding constraint on everything above."* Every calendar in the record derives from one or the other. **You must settle this.**
- **No CI**, in a product that wants to sell document CI. The other codebase has it. Porting is one day.
- **Four gates in this repo have reported green while blind.** That is not a process problem, it is a trust problem — our own tooling has lied to us.
- **The learning loop is not running:** 6,884 routing decisions in seven days, one reward label. 0.01%. It may not be described as self-improving.
- **The 41-day MVP estimate assumes full-time work.** With 54–78 hours a month of client work, real elapsed time is roughly **2.5× the estimate.**

PART III

What to build

7 The product, restated

GOOD

frontmatter is a markdown editor that knows which bytes a machine wrote. Every AI edit is recorded as a byte range with its prompt, model and timestamp. Machine-written spans look different from yours. One keystroke reverts any of them. Nothing else in the file moves.

Why this and not the other things.

- It is the **only** use of the engine that answers a complaint people are actually making — slop, 23.7% and rising 149% a year.
- It cannot be copied by a regenerating editor. Cursor and Zed rewrite whole files; they cannot tell you which bytes were theirs because they do not know.
- It reuses about **90%** of what already exists.
- It is a feature of an editor, which is the form factor you have chosen.

8 MVP-0 — the proof · 6 weeks

One demo: open a document an agent has been editing. Every machine-written span is visibly marked. Hover shows the prompt and the model. One key reverts it. `git diff` shows nothing else changed.

Ships	Why
NF-1 + NF-2 + NF-3 engine fixes	83% of foreign vaults are refused today. Five of six user flows start with this step
Provenance store — byte range, prompt, model, timestamp	The product
Provenance rendering — machine spans marked	The demo
One-key revert of any span	The proof
CI, with a deliberate red run first	Four gates have reported green while blind
Quick-switch, command palette, search	Table stakes. Without them nobody finishes the review

Not in it: sign-up, billing, sync, mobile, publishing, hosted AI, renders of any kind.

Exit criterion: ten strangers, their own repos. **Six of ten say they would keep it.** Not "like it" — keep it.

9 MVP-1 — the first money · +8 weeks

- **Vault-wide refactor** — rename a tag, a heading, a property key; every affected link shows as a reviewable hunk; refuse on ambiguity. This is the 86-like problem.
- **Free forever** for the editor. Charge for what the category pays for.
- **First paid line: teams.** Provenance across a shared repo, with a per-seat licence — the Obsidian commercial model, which is \$50/user/year and has no enterprise features at all.
- Billing, a support channel, ToS, and a way to tell users about a breaking change. None exist today.

10 MVP-2 — the moat · +12 weeks

- **Sync, done provably safely.** The #1 loved feature, the #1 switching trigger, and the only price this category has ever proven. A competitor's sync duplicates sections of files; ours cannot, and that is demonstrable.
- **The free wedge:** nested-construct live preview, shipped as an Obsidian plugin. 501 likes. It is the cheapest audience we will ever buy.

PART IV

The money

11 What it costs and what it pays

Line	Number	Basis
Monthly nut, two people	~₹1.09L	Kolkata cost of living + infra + tooling
Consulting hours to cover it	54–78/month	At your measured ₹1,400–2,000/hour
Hours left for product	~120/month	Out of ~198
MVP-0 at that rate	~10 weeks elapsed, not 6	2.5× the full-time estimate
Two-year success case	₹1,09,135/month	502 paying users, 171,200 cumulative visitors

TEST

Fund it with services, not a raise. ₹4 Cr is 367 months of the nut, from a fund with a seven-year horizon. That mismatch is the argument. Four fixed-scope engagements a year at ₹3,00,000 covers everything and costs roughly one week a month.

12 Before you write any code

Two weeks. ₹0. This is the highest-value fortnight available.

#	Test	Kill signal
1	Show 10 AI-heavy developers a mock: machine-written spans marked, one-key revert. Ask what they would pay	Fewer than 4 of 10 call it useful without prompting

#	Test	Kill signal
2	Post the live-preview fix as an Obsidian plugin	Under 200 installs in 14 days
3	Ask 5 people who bill for documents whether "which parts did the machine write" is a real problem	Nobody has been asked for it
4	Price test: three landing pages — free, \$8, \$20/seat	No email captures at any price

- If tests 1 and 3 both fail, **the provenance thesis is dead** and the answer is services plus a free plugin.
- If test 2 succeeds, you have an audience before you have a product. That has never been true before.

PART V

What our own review disproved

13 Corrections to the record

EXHIBIT 4

Our own review, against our own case

#	What we believed	The correction	Effect
C1	RCM starts after ₹20 lakh	CGST §24(iii): no turnover floor. Buying Claude API access is importing a service	GST registration starts at the first rupee, not at ₹20 lakh
C2	IT Rules set a 50 lakh threshold	The Rules carry no number — "as notified by the Central Government"	The threshold is not ours to plan against
C3	EU representative under GDPR Art. 27	DSA Art. 13 is a separate mandate with no small-enterprise exemption	Different obligation, different cost
C4	Vanta/Drata charge \$7,000–30,000	Neither publishes a price. drata.com/pricing returns 403	Our B2B ROI headline had no source
C5	GitBook's top complaint is lost work	Of 1,190 discussions, 51.3% are feature requests	Moat #2's only external evidence
C6	frontmatter and sgnk-md are siblings	202 of 228 files byte-identical. 0 files exist only in md	They are one codebase. Park md
C7	The corpus holds 7,959 frontmatter files	7,969. Counted four ways	Not definitional. Just wrong
C8	Our table editor is fine	It rewrites 3 of 4 lines and destroys 3 of 4 carriage returns [measured here]	Dead code. Keep the measurement as a test
C9	Ghost text is refused	It ships, defaulting to off , exactly as specified	No contradiction. An agent got this wrong; I checked

NOTE

62% of our load-bearing claims needed correction when someone opened a primary source: 13 confirmed, 13 revised, 5 refuted, 3 unverifiable. Re-derive before quoting anything in the record.

14 The decisions only you can make

#	Decision	Why it is first
1	One founder or two?	Every calendar and cost number depends on it. The record contradicts itself
2	Which buyer — the AI-native developer, or the person who is liable when a document is wrong?	They want different products. We have been building for both
3	Do you accept provenance as the product , and byte-exactness as the mechanism rather than the pitch?	It changes the demo, the copy and the price
4	Editor free, charge for teams and sync?	Inverts the revenue plan
5	The two unrotated PATs	Not a decision. An action. Today

RISK

What would make me tell you to stop. If the two-week tests come back with fewer than 4 of 10 developers calling provenance useful, and nobody who bills for documents has ever been asked for it, then there is no product here — only an engine, a consulting business, and a free plugin that makes people happy. That is not a failure. It is a smaller, truer version of the same work, and it pays better sooner.