

frontmatter

The decision layer — 20 minutes, and only this

This is the only document you need for the founding session. It carries no evidence of its own; every claim points down to the section that holds it. That rule is what keeps it short, and it is why you can trust a number here — it was measured somewhere you can go and check.

How to read it. Twenty minutes, cover to cover. The record underneath is 297,614 words and 125 sections; you do not need it to decide, only to verify. When a line here matters enough to argue about, open the section it cites and argue with that.

Contents

01 The finding that changes the plan

1.1 And the gap nobody is working on

02 The product, in one paragraph

03 What we ship

3.1 MVP-0 — the proof · 38 pts $\approx$ 41 days · 3.2 What the new evidence adds to the roadmap · 3.3 What to cut from the copy, today

04 The flows

05 Who buys it

06 Money

07 What is already decided, and may not be reopened

08 Open decisions, in dependency order

09 What would tell us we are wrong

10 Things that are true about us, not about the market

11 Where everything else lives

§1 The finding that changes the plan

Seven research rounds and ninety-two agents went out to test what we already believed. **Two of our three central beliefs came back refuted, from counted public evidence.** That is the most valuable outcome available, and it should be read before anything else in this document.

WHAT WE BELIEVED	WHAT THE EVIDENCE SAYS	WHERE
People will switch for byte-exactness	4 complaints about mangled formatting in 12,556 editor comments. One in 3,749 launch comments. One formatting-caused switch in eleven years. "line endings" appears in 1 of 43,656 GitHub issues	r24-v2
Generating handovers and context packs is the wedge	89 such products launched on Hacker News in 20 months. Median 2 points. 88 of 89 never reached 50. Four were named <i>Handoff</i>	r21-s1
Decision-flow renders are the obvious first render	The least-demanded render measurable anywhere — 563 downloads of 143,283,562 in the Obsidian registry (0.0004%), 3 plugins of 7,079, 0 of 1,600 Reddit posts	r21-s6

None of this kills the product. All of it changes the pitch.

The mechanism was never the problem. A byte-exact engine that refuses rather than guesses is still the right architecture — three days before we looked, **Cursor's own staff described doubled blank lines and destroyed carriage returns on CRLF files as "a known issue we're tracking"** when their agent does a full-file rewrite (r24-v3). The incumbents have the defect. We do not.

What changed is that "we don't corrupt your bytes" is a promise nobody is asking for. It has to be sold as the thing it enables, not as the thing it prevents.

1.1 And the gap nobody is working on

Slop — verbose, generic, unmaintainable AI output	23.7% of all complaints, growing +149% in 20 months
Classified	WORKFLOW, not model — durable, and ours to solve
Currently addressed by this product	Nothing

This is the fastest-growing complaint in the entire corpus and the product does not touch it (r21-s1). It is also the one thing the artefacts plausibly answer: a decision record, a spec and a flow are not valuable as documents — they are valuable as **the context that makes the next AI output reviewable instead of slop**. That is a different product story from "we generate your handovers", and it is the one the evidence supports.

§2 The product, in one paragraph

frontmatter is a markdown editor for people who build software with AI. It opens a folder of .md files — in a GitHub repo or on disk — and edits them byte-exactly, changing only the bytes you touched. When an agent proposes a change, you see the exact diff and accept or reject it hunk by hunk, and what you rejected is bit-identical afterwards.

One line for someone who will not read that one: **it turns your repo into the prompt** (§104).

What it is not: not Notion, not a project manager, not a chat wrapper, not an IDE. Those are refusals, not omissions (§103).

§3 What we ship

3.1 MVP-0 — the proof · 38 pts ≈ 41 days

The demo is a rejection. Point an agent at a real repository. Let it make a six-file change. Reject one hunk. Run `git diff` and show the rejected file is byte-identical.

It is the only demo where the absence of an event is the product (§95.2).

SHIPS	PTS	WHY IT IS IN
NF-3 — bare-CR fence, set-destructive	4	A correctness demo cannot ship on a writer with a known silent-destruction path
NF-1 + NF-2 — zero-indent sequence	6	83% foreign-vault refusal. A stranger's repo <i>is</i> a foreign vault
CI, with one deliberate red run first	2	Four gates in this repo have already reported green while blind
Real byte budget, replacing the stub	2	Currently a literal <code>echo</code> in a product that wants to sell document CI
Seams 2 + 3 — one splice behind every write	8	Today one symbol from one of thirteen engine files reaches product code
<code>Land()</code> + per-hunk accept/reject	16	The demo itself

Deliberately absent, and say so on the page: no sign-up, no hosted anything, no billing, no sync, no mobile, no kanban, no calendar, no publishing.

Exit criterion, observable not felt: ten people who are not friends, family or in our Discord, each given the binary and their own repository. Six of ten run it on a real repo.

The concession we make out loud: for one person editing their own prose, there is no reason to use this over a text editor and git. We should say that rather than pretend otherwise (§95.2).

3.2 What the new evidence adds to the roadmap

These are not in MVP-0 and they are the strongest candidates for MVP-1, because unlike byte-exactness they are things people are complaining about today:

BUILD	EVIDENCE	§
Vault-wide refactor — rename a tag, a property key, a link target across the vault, with a preview diff and a refusal when a target is ambiguous	86 likes on broken-links-on-rename. This is byte-exactness sold as a capability rather than a promise	<code>r24-v2</code>
Nested-construct live preview	501 likes — the single most-voted bug in Obsidian's history — plus 96 + 82 + 50 + 36 + 33 on siblings. Our OffsetMap makes it tractable	<code>r24-v2</code>
Sync, reframed as "provably safe on top of whatever you already use"	#1 loved, #3 hated, #1 switching trigger, 1,251 of 43,656 issues. Two top bugs are a competitor's sync duplicating sections of files	<code>r24-v2</code>

3.3 What to cut from the copy, today

- The markdown-flavour argument.** 17 of 3,749 launch comments (0.5%) mention CommonMark or GFM at all. It is an engineering constraint, not a message (`r24-v2`).
- "Protection against lost work."** First-hand destruction is ~0.05% of comments (`r21-s1`).
- Any roadmap item aimed at hallucination, mid-session incoherence, over-eagerness or wrong-file edits.** Measured decline of 19–67% in 20 months — the labs are fixing these (`r21-s1`).

§4 The flows

Five of the six user flows begin with the same step, **and that step currently fails 83% of the time** — it is the NF-1 zero-indent refusal (r21-s7). That single defect gates the product's entire front door, which is why it sits in MVP-0 rather than in a backlog.

The handover flow, which the record was proudest of, should be **demoted from a button to an outcome** (r21-s7). Nobody wants to press "generate handover". They want the next session to go well.

§5 Who buys it

The B2B finding is counter-intuitive and worth the whole section: the closest structural analogue to this product runs a nine-figure-logo B2B business **with zero B2B features** — the entire enterprise offering is a \$50/user/year licence (r21-s5).

The implication is that our B2B strategy is not SSO, SCIM and an audit trail. It is a commercial licence and the fact that documents already live in the customer's own repository, which means data residency and exit are solved by the architecture rather than by a roadmap.

D2C is the harder half: developer free-to-paid conversion is the hard half of freemium, and we have zero paying users. That is a fact about us, not a market condition.

§6 Money

Services fund this. Not investors, not the product. (r22-t3)

Client hours covering the whole two-person nut	54–78 per month
A ₹4 Cr raise, expressed in that nut	367 months of runway
From a fund with a horizon of	~7 years

That mismatch is the argument against raising, stated in its own terms: thirty years of runway bought with seven years of obligation.

The contradiction you must settle before any of these numbers mean anything: this document says two founders. §53 says *"One person. This is the binding constraint on everything above."* Every capacity, calendar and cost figure in the record derives from one or the other. They cannot both hold.

§7 What is already decided, and may not be reopened

DECISION	SETTLED BY
frontmatter is sgnk-md — 202 of 228 shared source files byte-identical, 0 files exist only in <code>md</code> . One codebase; <code>md</code> is parked	Measured, <code>r22-t2</code>
Nothing else in the ecosystem fuses. HQ, CareerOS, Markex, Advox, Brand OS, INW, skills-registry, stock, trade — all NO. Most share a founder and nothing else	<code>r22-t2</code>
AIOS does not ship as a product , and §15's "yes" is deferred, not honoured	<code>r22-t2</code> , <code>r23-u3</code>
Sync is git merge + splice journal + compare-and-swap. Never a CRDT	§31.1
Documents live in the user's repo and never move	§5
No arbitrary client-side code execution, ever	§9.2

§8 Open decisions, in dependency order

Ordered by what each unblocks, not by how interesting it is.

#	QUESTION	WHY IT IS FIRST	UNBLOCKS
1	One founder or two?	Every capacity number in the record depends on it	The entire calendar
2	Do we accept the repositioning — mechanism not pitch, refactor not byte-purity?	MVP-0's demo assumes the old pitch	The demo, the copy, the launch
3	Offline-first desktop, or web?	Determines file access, the legal surface, and where AI keys live	§91, §100, §101
4	BYO key or platform key?	Sets the free tier's arithmetic limit	§90, §93
5	Does the agent get write authority?	The flagship demo and the liability posture	§14, §100
6	Publishing in v1?	A permanent, personal, unbounded on-call obligation	§44
7	CJK in v1?	Engine scope	§33
8	The two unrotated PATs	Only you. Today. Not a decision — an action	§54 D5

§9 What would tell us we are wrong

Each of these is observable, dated, and cheap enough to check before committing the next quarter.

BET	FALSIFIED WHEN	BY
Byte-exactness is a felt need	Fewer than 3 of 10 strangers care, given the binary and their own repo	Day 41
The 83% refusal is one bug, not a class	First patched corpus run still refuses >10 of 7,969	Day 10–14
Our own gates can see	CI passes on a deliberately broken commit	Day 2
Someone pays for hosted inference over a file they own	60 days, working checkout, zero non-founder paid signups	Day 130
Free is not too generous	Free-to-paid under 2% at $n \geq 200$ free users	Day 160
The certificate is a standard	90 days after publishing the spec, zero third-party runs	Day 220

§10 Things that are true about us, not about the market

Stated plainly because they will otherwise be discovered at a worse moment.

- **62% of the record's load-bearing claims needed correction when someone opened a primary source.** 13 confirmed, 13 revised, 5 refuted, 3 unverifiable (§89). Re-derive before quoting.
- **Reverse charge has no turnover floor.** CGST §24(iii) compels GST registration for anyone liable under RCM. Buying Claude API access is importing a service. Registration starts at the first rupee, not at ₹20 lakh (§89.3).
- **The learning loop is not running.** 6,884 routing decisions in seven days, one reward label — 0.01%. The instrumentation is alive; the learning is not (`r23-u1`). "Self-improving" may not be claimed.
- **We have no CI**, in a product that wants to sell document CI. `md` has it; frontmatter does not. Port it first, ~1 day (`r22-t2`).
- **Our own table editor destroys bytes** — editing one cell in a 4-line CRLF table rewrites 3 lines and destroys 3 carriage returns [measured here]. It is dead code and already slated for deletion. Delete the UI, **keep the measurement as a test**: it is a working demonstration, in our own repo, of exactly what regenerating costs.
- **The documentation outgrew the product.** 2,219,390 words of markdown against 25,407 lines of source — 87 words per line [measured]. This document exists because of that ratio.

§11 Where everything else lives

`docs/MAP.md` routes the rest. The short version:

QUESTION	OPEN
What is this, in one page	THESIS.md §104
Every feature ever proposed, and the ones refused	PRODUCT.md §94
Where the MVP line falls	PRODUCT.md §95
How to attack all of it	PRODUCT.md §96
Is this number trustworthy?	VERIFICATION.md §89, always
What the engine is and what is wrong with it	ENGINE.md §67–80
The stack	DEV-PLAN.md
The raw research	docs/research/ — open it to check a claim, never to find an answer