

The pilot plan — what we build first, what we cut, and how we sell it

Amit — the plan for the meeting. One decision on the first page, the market in one table, every component with a verdict, three tiers, the screens as they will actually look, the price, the channels, and ninety days. Nothing is said twice. Every number was fetched from a primary source this week, and the ones that failed a second check were removed.

KEY FIGURES

496,000

stars on the free tools that already generate specs and plans from an idea

0

of them is an editor. That is the product

6

AI routes already live in our repo, wired to a five-provider fallback chain

83%

of real vaults refused today by one bug. Four days to fix

114

paying seats at \$8 covers the whole monthly nut

PART I

The decision

1 What we build

GOOD

frontmatter is the editor where your agent's specs and decisions get read, reviewed, and finished. Open a spec the agent wrote. Everything it wrote is visibly marked until a person has read it; one key reverts any of it; and the questions it left open — [NEEDS CLARIFICATION: auth method - email, SS0, OAuth?] — are answered in place, one byte range replaced, nothing else in the file moved.

One honest caveat, found while checking this page. spec-kit's own /speckit.clarify is chat-first: it asks the questions in the agent, one at a time, and writes the answers back itself. So answering is not the front door — reviewing is. What makes answer-in-place real anyway: 62,976 public files under specs/ paths carry an unanswered NEEDS CLARIFICATION today (GitHub code search; the phrase is generic and template copies inflate it, so treat it as an upper bound). Questions survive into committed files. Someone has to find and close them, and nothing currently shows where they are.

The proposed pilot was: idea → decisions → generate every file → kickoff prompt → your agent executes. That flow already ships, free, from GitHub and three others (§2). What none of them has is a surface: they write markdown and hand it to whatever editor you happen to use, and that editor has no idea an agent wrote it. So we are not the generator. We are where the generated thing lives.

Why this is the first framing that is actually ours. Answering a question inside a file without disturbing the rest of it is a byte-range splice — the one thing our engine does and a whole-file rewriter cannot. It aims at slop, the fastest-growing complaint in the market (23.7%, +149% in 20 months) and the only one nobody is working on. And it turns half a million stars of free tooling from a competitor into a channel, which matters because every other channel we have is borrowed.

The fifteen-second demo. Run /speckit.specify in your agent. Open the spec it wrote, here: every paragraph is tinted, because nobody has read it yet. Hover one — who wrote it, when, the prompt. Revert one with a key; git diff shows only that. Then the panel on the right lists the two questions the agent left open; answer one inline, and git diff shows one line.

PART II

The market, in one page

2 What exists, what it costs, and where it stops

EXHIBIT 1

Who already does what, fetched 2026-09-06

Tool	Reach · price	What it does	Where it stops
Generators obra/superpowers	282,242 stars — #13 on all of GitHub · free	Brainstorm → design → writes the spec to docs/.../specs/, commits it, review gate	Terminal. Not an editor

	Tool	Reach · price	What it does	Where it stops
	github/spec-kit	133,660 stars — only 73 repos have more · free	/speckit.specify · /speckit.clarify · /speckit.plan · /speckit.tasks · /speckit.implement · a go/kill assess extension	Writes markdown, then leaves. No provenance
	Kiro (AWS)	GA · \$0 tier	requirements.md · design.md · tasks.md, with ambiguity detection	Its own IDE, code-first
	BMAD · task-master	52,719 · 28,053 · free	Planning and task breakdown inside the agent	Same
Editors	Obsidian	Editor free · Sync \$4 · Publish \$8/site · \$50/user/yr commercial	The category leader	No AI editing worth the name
	Cursor · Zed	~\$20 · free	The strongest agent editing for code; Zed reviews per hunk	Whole-file rewrite. Cursor staff admit destroying CRLF. No markdown vault semantics
	Notion	~\$10/user	Teams, databases	Files are not files. Export is named blocksToMarkdownLossy
Builders	Lovable · v0 · bolt	\$100M+ ARR published · \$25–30/mo	Idea → app, for people without an agent	A different buyer entirely. Their ARR does not validate us
Browser	Obsidian Web Clipper · Markdown Viewer	1,000,000 · 500,000 Chrome users	Capture and read markdown in the browser	Reading only. A channel, not a competitor
Hosted sites	mdown.ai	Hosted, aimed at non-engineers	Markdown → live website, on their servers	The host role we refuse

What the price column settles. The editor is worth ₹0 in this category. The two things people have proven they pay for are **sync (~\$4)** and a **commercial team licence (~\$50/user/year)**. Our buyer already spends \$20–40 a month on AI tooling, so we ask for a share of an existing budget, not a new line.

3 What was AI jargon in the original plan, and what survives

EXHIBIT 2 The audit of our own vocabulary

It said	What it actually was	Verdict
"Context packs" and "handover generation"	A category with 89 launches and one hit. And now free from GitHub	Cut
"Decision-flow renders"	A note-taker's plugin at 0.30× the median. Wrong buyer	Cut
"Kickoff prompt", "reads our CDN files"	A copy-paste step worse than a slash command; a hop that contradicts <i>we never hold your documents</i>	Cut
"Self-improving AIOS"	6,884 routing decisions produced one feedback label	Cut from the pitch. Internal tooling stays internal
"Local LLM fallback"	Five <i>cloud</i> free-tier providers, not a local model. And a constrained small model fabricates rather than refuses (fields drive fabrication to 100% in 10 of 13 models)	Defer. Keep the cloud chain for testing
"Byte-exactness" as the headline	4 complaints in 12,556. Nobody asks for it by name	Keep as the mechanism, never the pitch
"Degradation certificate", "byte-level time travel", "projection law"	Engineering vocabulary	Internal only
Provenance — see what the machine wrote, undo any of it	Answers slop. Uncopyable by anything that rewrites files	The product
Answer in place	A byte splice into a file 496k stars of tooling already produce	The pilot's point
File-scoped AI	Table stakes — 5 of 6 incumbents ship it	Keep, bring-your-own-key

What we ship

4 Keep, cut, defer

EXHIBIT 3

Every proposed component, decided

#	Proposed	Verdict	Why
1	The markdown editor — Live · Edit · Split · Read	KEEP	The free tier. Zed's launch page had zero AI words in 6,851 characters; Obsidian's first build had no AI. Editors ship the editor first
2	The design system	KEEP	Written down already. The shipped <code>globals.css</code> is a sibling project's — replace it
3	Decisions generated from the idea	RESHAPE	Demand is real, +40% year on year — and <code>/speckit.clarify</code> , Kiro and superpowers ship it free. We answer; we do not ask
4	Generate every project file	CUT	<code>/speckit.plan + /speckit.tasks</code> , free, inside the agent. Our idea mode (one document from one idea, five kinds) already ships — keep that
5	Kickoff prompt	CUT	The slash commands are the kickoff prompt
6	Agent reads files from our CDN	CUT	Buys nothing (<code>spec-kit</code> writes to the repo, where every agent looks) · Codex ships with network off · context7 at 61,691 stars refuses this pattern and ships MCP instead · breaks our one promise
7	AI scoped to folder / file	RESHAPE	Keep, file-first . A folder is a pointer, not a content dump — Claude Code itself reduces a directory to "file listings, not contents"
8	Local model	DEFER	See §3. The cloud chain is right for pre-launch

5 The three tiers

EXHIBIT 4

MVP-0 pilot · MVP-1 pro · MVP-2 max

	MVP-0 · Pilot — 10 weeks, free forever	MVP-1 · Pro — teams, \$8/user/mo	MVP-2 · Max — +\$4/user/mo
Engine	NF-1 + NF-3 fixed · CI with one deliberate red run · engine wired behind every write · real byte budget	Vault-wide refactor with a reviewable diff, ambiguity refused	Sync, provably safe — git merge, never a CRDT; the #1 loved feature and the one proven price
Editor	Four modes · tree with the unreviewed bar · tabs · quick-switch · palette · search · properties panel	Conflict view	Multi-device
Provenance	Store (range, prompt, model, time, in a sidecar in <i>their</i> repo) · tinted spans · hover card · one-key revert · review panel	Across a shared repo: who reviewed what, per person · commercial licence	Team admin, seats, roles
Answer in place	[NEEDS CLARIFICATION] and Q: → A: markers answered inline, one line changed	Questions panel across the whole repo	—
AI	File-scoped panel, bring-your-own-key only · idea mode as it ships today · structured output in the port	Hosted AI, metered, hard-capped, never default	Document gates in CI: link integrity, stale sections, unreviewed machine text
Output	—	Site · page · deck from the folder. We never host it	Push to Vercel / GitHub Pages / Netlify, their account
Channels	Free Obsidian plugin, week 2 (the 501-like live-preview fix)	MCP server so their agent writes into frontmatter directly · Chrome extension that opens any <code>.md</code> URL here	—
Deliberately absent	Sign-up · billing · sync · mobile · hosting · hosted AI · local model	Mobile · plugins · a chat sidebar	Plugins, still. Mobile as a reader only
Exit criterion	6 of 10 strangers keep it after two weeks	A team pays within 20 qualified conversations	Sync passes the corpus byte-identical

Before any stranger touches the pilot, three things that do not exist today must: bring-your-own-key (settings expose one boolean; every request bills us), structured output in the LLM port (it returns a bare string), and CI (four of our gates have reported green while blind).

6 How it is built — the plugin question, answered

The most cost-efficient shape is the one where **we run no model, hold no document, and sit inside the tools the user already pays for.**

EXHIBIT 5

The delivery surfaces, and what each is for

Surface	Role	Cost to us	When
Web app	The trial. Visit a URL, connect a repo, nothing to install	~\$0	MVP-0
Desktop (Tauri)	The daily surface. Local folders, offline everything except language AI. Same build as the web app — one filesystem adapter apart, so we never maintain two products	\$99/yr Apple notarisation; Windows signing \$150–400/yr with a hardware token	MVP-0
Their agent, via MCP	The AI runs in Claude Code / Cursor / Codex on their subscription. frontmatter is the surface the agent writes into and the record of what it wrote. Zero inference cost, zero key handling, and the agent they already trust	~\$0	MVP-1
Bring-your-own-key panel	For AI inside the editor. Their key, their provider, their machine → provider. Never through us	~\$0	MVP-0
Obsidian plugin	Distribution only. The live-preview fix in their store, permanently	4 days	Week 2
Chrome extension	Distribution. Open any .md on GitHub in frontmatter, provenance shown. 1.5M users already read markdown in the browser	~1 week	MVP-1
Our hosted AI	Optional, metered, hard-capped, never the default	Pass-through	MVP-1
Vendor CDN · hosting	—	—	Never

So the answer to "do we ask them to install a plugin": no. The plugin is a channel. The product is the web app and the desktop build, and the AI is theirs — through MCP or their own key — which is why the pilot's AI cost is zero and its trust story is "the request never touches us."

The rules that keep the architecture cheap and honest: documents never enter our control plane, which holds identity, teams and billing and zero document bytes · nothing runs in the background, so spend follows use rather than time · **cap output, not input** — output is 75.8% of spend · the byte↔UTF-16 seam is the highest-risk thing we own, and anything crossing it without the mapping layer corrupts non-English text silently · provenance is a plain-text sidecar in *their* repo, readable without us.

PART IV

The screens

7 As they will actually look

Built from the two hand-drawn layouts, with real content. **Clickable prototype:** claude.ai/code/artifact/f6baec30-9d0f-4e4f-a231-50f178f0e65e — also in the repo at docs/prototype/frontmatter-prototype.html. Every screen is reachable from the strip at the bottom.

EXHIBIT 6

S0 - The launcher — blank, four templates, existing edits sorted unreviewed-first

Fm

Search every document, every project

Open folder

New

START SOMETHING

Blank document

empty file · yours from byte one

Spec

what you are building, and what you are not

Decision record

MADR · context, options, consequences

Handover

what the next person needs

Changelog

Keep a Changelog 1.1.0

EXISTING EDITS · 3 PROJECTS · 128 DOCUMENTS

unreviewed first

auth.md

acme-billing / specs

GitHub App installation flow rather than an OAuth app, because the App grants permission per repository...

64% unreviewed

agent · 2h ago

0004-sync.md

acme-billing / adr

We will not use a CRDT. Merging by git three-way keeps CRLF and the missing final newline...

88% unreviewed

agent · yesterday

pricing.md

acme-billing / notes

Team \$8/user/month. Sync \$4 later. The editor is free because the category settled that.

31% unreviewed

you · 3d ago

billing.md

acme-billing / specs

One payment attempt per Indian card. ₹15,000 per transaction is an architectural constant.

12% unreviewed

agent · 4d ago

README.md

acme-billing

Billing for small teams. Runs on your repo; we never hold your documents.

all reviewed

you · 1w ago

0003-scope.md

acme-billing / adr

Out of scope for v1: invoices in currencies other than INR and USD, and any refund automation.

45% unreviewed

agent · 1w ago

onboarding.md

percent / notes

Minute zero: one button, choose a folder. No account, no email field, no tour.

all reviewed

you · 2w ago

api.md

percent / specs

REST with idempotency keys on every write. Errors are RFC 9457 problem details.

22% unreviewed


agent · 2w ago

3 projects · 128 documents · 7 unreviewed

⌘K commands · ⌘P files

The bar on every card is the product showing itself before anything is open. Templates are markdown skeletons, not a gallery: spec, decision record, handover, changelog — the four document types the research found actually recur.

S2 · The editor — tree, coloured tabs, four modes, provenance, and the AI section on load



Machine-written text that nobody has read carries a light tint and nothing else — no border, no icon, so it can be ignored. Click a span for the card: who wrote it, when, which bytes, the prompt, **mark reviewed** or **revert**. The AI section at the bottom is the Google-Docs-style writing panel from the sketch: it shows on load, it proposes marked spans, it writes nothing until you keep it, and it runs on **your** key.

EXHIBIT 8

S3 · Answer in place — a spec-kit spec, three questions, one answered

fm

spec.mdplan.mdtasks.md

ShareSearchsm

acme-billing

specs

001-invoicing

spec.mdplan.mdtasks.md

.specify

AGENTS.md

LiveEditSplitRead

written by spec-kit via Claude Code · 11 min ago · 3 questions for you

feature: 001-invoicing

status: draft

generated-by: speckit.specify · 2026-09-06

Feature Specification: Invoicing

Functional requirements

- FR-001: System MUST let an admin create an invoice from a customer record with line items, tax and a due date.
- FR-002: System MUST render the invoice as a PDF that is byte-identical across renders for the same input.
- FR-006: System MUST authenticate users via
 - A SSO through the GitHub App – no second identity path
- FR-007: System MUST retain invoices for
 - NEEDS CLARIFICATION 7 years – Indian Companies . Answer
- FR-009: System MUST charge Indian cards
 - NEEDS CLARIFICATION once, or retry on failure?

Non-goals

Multi-currency invoicing, refunds automation, and any accounting export are out of scope for this feature.

Answering writes → A: _ into the file at that byte range, exactly where spec-kit reads it back. 1 line changed, 0 lines moved. The requirements the agent wrote stay tinted until you have read them.

Questions in this spec2

FR-007 · retention periodopen

FR-009 · card retry policyopen

FR-006 · auth methodanswered

Document outline

Functional requirements5 · 5unreviewed

Non-goals1 · 1

Document history

nowyou · answered FR-006

11mspec-kit · /speckit.specify · 2,140 bytes

Mark all readShortcuts

AI edit · #J

7 unreviewed spans · 92%2 questions open

git · +1 -1 · spec.mdLF · UTF-8

The moment the pilot exists for. spec-kit wrote the file eleven minutes ago; FR-006 is answered, FR-007 is being typed. The status bar reads git · +1 -1 — one line, nothing else moved.

EXHIBIT 9

S4 · The review panel — catching up on a file you did not watch being written

0004-sync.mdauth.mdpricing.md

ShareSearchsm

acme-billing

specs

auth.md

adr

0003-scope.md

0004-sync.md

notes

pricing.md

LiveEditSplitRead

acme-billing / adr / 0004-sync.md

title: ADR-0004 Sync without a CRDT

status: proposed

date: 2026-09-02

ADR-0004 · Sync without a CRDT

Context

Two devices edit the same markdown file. Most sync products reach for a CRDT because convergence is guaranteed. Convergence is not the property we want: CRDTs interleave concurrent edits character by character, so two people finishing the same sentence produce a merged sentence neither wrote.

Decision

Git three-way merge, with divergence surfaced as a reviewable conflict. On the same inputs it preserves Windows line endings and a missing final newline, and it stops exactly where a careful writer would want to stop.

Consequences

Conflicts happen, and users hate conflicts. We publish the conflict rate as a shipped metric with a budget.

We never guess which version was meant. That is the whole point.

3 UNREVIEWED IN THIS FILE

Context · 312 bytes · sonnet-5 · 2h

Two devices edit the same markdown file. Most sync products reach for a CRDT because...

a · r · s ·

revert · skip

Decision · 218 bytes · sonnet-5 · 2h

Git three-way merge, with divergence surfaced as a reviewable conflict...

Consequences · 96 bytes · sonnet-5 · 2h

Conflicts happen, and users hate conflicts. We publish the conflict rate...

Provenance is a state, not a moment. These stay marked until a person accepts or reverts them — today, next week, or when a teammate opens the file.

3 unreviewed spans · 88% · 231 words

git · clean · main · CRLF · UTF-8

Every unreviewed span in order, j k to move, a r s to accept, revert or skip, and a counter that reaches zero. Provenance is a state, not a moment: these stay marked until a person deals with them.

EXHIBIT 10

S2 · Split — source and rendered, the same bytes

fm

pricing.md x auth.md x

Share Search sm

< TREE FP+

acme-billing

- notes
- pricing.md
- specs
- adr

Live Edit Split Read

acme-billing / notes / pricing.md · source = rendered, same bytes

title: Pricing

status: decided

tags: [pricing, obsidian, teams]

Pricing

The editor is free. Obsidian settled that: the whole editor at \$0, sync at \$4, a \$50/user/year commercial licence. Everyone priced around it.

What we charge for

- **Team** - \$8/user/month. Provenance across a shared repo, review state across people, admin.

- **Sync** - \$4/user/month, later, when it is provably safe.

The arithmetic

At \$8/seat, **114 paying seats** covers the monthly nut of ₹1.09L - about twenty small teams. Churn on a team licence is far lower than on a personal one.

title: Pricing

status: decided

tags: [pricing, obsidian, teams]

Pricing

The editor is free. Obsidian settled that: the whole editor at \$0, sync at \$4, a \$50/user/year commercial licence. Everyone priced around it.

What we charge for

- **Team** — \$8/user/month. Provenance across a shared repo, review state across people, admin.
- **Sync** — \$4/user/month, later, when it is provably safe.

The arithmetic

At \$8/seat, **114 paying seats** covers the monthly nut of ₹1.09L — about twenty small teams. Churn on a team licence is far lower than on a personal one.

2 unreviewed spans · 31% 96 words

git · clean · main LF · UTF-8

EXHIBIT 11

S11 · Settings — one page, and the key field that does not exist yet

fm Settings x sm

Settings

APPEARANCE

Theme

System

Editor typeface

Google Sans Code

KEYMAP

Default

AI PROVIDER

Provider

Anthropic

Your API key

.....

Stored in your OS keychain. Requests go from this machine to the provider. They do not pass through us.

Model for language work

claude-sonnet-5

Structural work — links, slugs, frontmatter fills — never calls a model.

☐ Use frontmatter's hosted AI instead (metered, hard-capped, never the default)

PROVENANCE

☒ Mark machine-written text until a person reviews it

☒ Show the unreviewed bar in the file tree

Record

byte range prompt model time session id

Written to .frontmatter/provenance.jsonl beside your documents, in your repo. Plain text. Readable without us.

GIT IDENTITY

Commit as

EXHIBIT 12

S8 · Refactor preview (MVP-1) — the engine made visible

fm

Rename tag #auth → #authentication

✕

8 FILES · 6 CHANGE · 2 REFUSED

specs/auth.md2 hunks

specs/billing.md1

specs/api.md1

adr/0003-scope.md1

notes/pricing.md1

README.md1

notes/2026-08-12.mdrefused

archive/old-auth.mdrefused

specs/auth.md · frontmatter · bytes 41-451 of 2

4 tags: [auth, github-app, security]

4 tags: [authentication, github-app, security]

accept

skip

accept all in file

specs/auth.md · body · bytes 1,902-1,9072 of 2

38 See also [[billing]] and the

39 #auth notes from the kickoff.

39 #authentication notes from the kickoff.

accept

skip

REFUSED · 2 FILES

notes/2026-08-12.md — #auth appears inside a fenced code block at line 17; rewriting code is a guess we will not make. archive/old-auth.md — the tag sits in flow form tags: auth with no brackets; the target is ambiguous. Both files are untouched.

Every file that will change is listed before a byte moves. Accept all, per file, or per hunk. Ambiguity is refused, never guessed — and the refusal says why.

Eight files, six change, two are **refused** with the reason — a tag inside a code fence, an ambiguous flow-form key. Nothing moves until you accept.

EXHIBIT 13

S12 · Generate (MVP-1) — a site from the folder, never hosted by us

fmGenerate · acme-billing ×

Write to ./site

Deploy to Vercel

OUTPUT

Static site · tree as navigation

One shareable page

Slide deck · H2 = slide

WHAT GOES IN

specs/ · adrl · README.md

notes/ excluded

Skip documents with unreviewed machine text

THEME

Plain · one column

Docs · sidebar

Anything custom is site, css in your repo.

DEPLOY

Vercel · your account

GitHub Pages

Netlify

We never host it. Files land in ./site in your repo. Deploy runs on your account, under your terms.

acme-billing.vercel.app/specs/auth

Specs

Authentication

Billing

API

Decisions

0003 Scope

0004 Sync

Authentication

We use the GitHub App installation flow rather than an OAuth app, because the App grants permission per repository instead of across the whole account.

The installation token is scoped to the repositories the user selected and expires after one hour, which means a leaked token has a bounded blast radius.

What we never do. Ask for the repo scope. Keep a second identity path.

generated by frontmatter · every paragraph reviewed by a person · source: specs/auth.md

A site is another projection of the same files. No site editor, no theme builder — edit the markdown, regenerate.

PART V

Money, market, and the plan

8 Price

EXHIBIT 14

What we charge, and what it covers

Tier	Price	What it is
Individual	Free, forever, no limits	The whole editor, provenance, answer-in-place, BYO key. This <i>is</i> the distribution strategy
Pro · Team	\$8/user/month	Provenance across a shared repo, who-reviewed-what, admin, commercial licence, MCP, generation
Max · Sync	+\$4/user/month	Multi-device sync, provably safe. Document gates in CI
Hosted AI	Metered, at cost plus a margin	Optional, never the default

The arithmetic: monthly nut ~₹1.09L for the two of us plus infrastructure, covered by 54–78 consulting hours; infrastructure is ~\$0 at 100 users and ~\$390/month at 10,000, because we store no documents. **114 paying seats at \$8 covers everything** — about twenty small teams. Funded by four fixed-scope engagements a year at ₹3,00,000, not a raise. **The real cost at scale is support, ~90 founder-hours a month at 10,000 users**, which is why the free tier carries no service commitment and we say so.

TEST

We give away the differentiator, deliberately. Provenance free is what makes anyone try it. The paid thing is not the feature — it is the feature *across people*. Downgrading loses features, never files.

9 What grabs users, and how we reach them

The message. Five words: *See what the AI wrote*. One line: *A markdown editor that shows which parts of your document a machine wrote — and lets you undo any of them*. For a team: *Know what nobody has reviewed yet, across the whole repo*. The proof line: *git diff* after an edit shows your change and nothing else. Tested on 8,513 real files, every release.

Lead with the refusals, not the features — this audience is sceptical of AI tooling: no plugins, no code execution, no holding your files, no lock-in, and if we disappear you lose nothing because nothing was ever taken. Never say *revolutionary*, *seamless*, *AI-native*, or *byte-preserving* in a headline.

EXHIBIT 15

Channels, in the order we use them

When	Channel	Why it is in this order
Week 2	Free Obsidian plugin — the live-preview fix	501 likes on the bug. The cheapest audience we will ever buy. Installs measurable by day 14
Week 2 →	The spec-driven toolchain	spec-kit and superpowers are MIT and agent-agnostic. A README line, an MCP entry, and "open the spec in frontmatter" in their docs turns 496k stars into an input
Now →	Our own writing	The research behind this document is the marketing. Nobody has published it
MVP-0 exit	Show HN	After 6 of 10 strangers keep it. One shot. Zed's scored 43, Cursor's 9 — launch points are near worthless as a signal
MVP-1	Chrome extension · Product Hunt	The extension reaches 1.5M people who already read markdown in a browser. PH is backlinks and a spike, not a strategy
Always	Generated pages with an honest mark	The only fully owned channel. A shared artefact is a feature and a channel at once

How we use the users we get. The ten pilot strangers are the research we have never done — zero user interviews to date. Every accept and revert is a preference signal we log. The plugin's installers are the list we do not have. And every generated page carries a small mark, which is the one loop that compounds without us.

10 Challenges, and what we do about each

EXHIBIT 16

Ranked by what would actually stop us

Challenge	What we do
Provenance turns out to be a nice-to-have	Two-week test before code. Kill gate: fewer than 4 of 10 call it useful unprompted
spec-kit or Kiro ship a review surface	Go deeper where they cannot follow: signing, export, an audit trail. Our provenance record is designed so it <i>can</i> become that without a rewrite
Zed or Cursor add markdown vault semantics — one sprint	Depth in markdown, not speed. They still cannot do provenance without abandoning whole-file rewriting
Distribution	The plugin, the toolchain, the extension — three channels that are cheap and two of which we did not have last month
The 83% refusal is a class, not a bug	Falsified if a patched corpus still refuses over 10 of 7,969, by day 24
Every AI call bills us	BYO key before any stranger; MCP so most AI never touches us at all
Client work crowds out product	Capped at 78 hours a month between us, written down
Two founders, no on-call	Free tier has no SLA and the page says so. One email address, a published response window
GST	Reverse charge has no turnover floor — registration starts with the first API purchase

11 Ninety days, and the two decisions

EXHIBIT 17

An observable outcome every fortnight

Weeks	What happens	Outcome
0	Rotate the two access tokens · decide the buyer	Done, written down
1–2	Four tests, build nothing : 10 developers shown a provenance mock · plugin shipped · 5 people who bill for documents asked · 3 landing pages at three prices	Kill signals read at day 14
3	Go / no-go	A written decision either way
3–4	CI with a deliberate red run · NF-3, then NF-1	CI fails a broken commit. Refusals fall from 83%

Weeks	What happens	Outcome
5–7	Provenance store and rendering · BYO key · structured output	A document shows machine spans on a stranger's key
8–9	Revert · answer-in-place · engine wired behind every write	The fifteen-second demo works end to end
10–12	Table stakes · ten strangers, their own repos and agents	6 of 10 keep it

EXHIBIT 18

The decisions

#	Decision	By
1	Accept the reshape — the editor <i>for</i> spec-driven work, not the generator <i>of</i> it. Everything above follows from this	The meeting
2	Which buyer first — the AI-native developer (reachable, pays little) or the person liable when a document is wrong (pays properly, hard to reach). My recommendation: the developer, with the provenance record built so it can become the liable buyer's audit trail later	Week 1
3	The two unrotated access tokens — an action, not a decision	Today

RISK

What would make me say stop. Fewer than 4 of 10 developers calling provenance useful, and nobody who bills for documents having ever been asked for it. Then there is no product here — only an engine, a services business, and a free plugin people like. That is not a failure; it is a smaller, truer version of the same work, and it pays sooner.

12 What I think

- **The engine is the best thing either of us has built**, and the funded competitors have admitted the defect it fixes.
- **We were selling it wrong twice** — as a promise about bytes nobody asked for, then as a generator competing with half a million stars of free tooling.
- **The third framing is the first one that is ours.** Everyone can generate a spec; nobody has a good place to read and answer one. Surfaces are sticky.
- **The two weeks of testing matter more than the ten weeks of building.**
- **Distribution is still the thing I am least sure of** — but for the first time we have three cheap channels and one we would own.