

The pilot plan — what we build first, what we cut, and how we sell it

Amit — the plan for the meeting. Re-checked 2026-09-06 against 157 primary sources; 60 corrections, three of them to the previous version's own page one, each marked where it lands.

KEY FIGURES

≈75–85k

public repos hold superpowers specs and plans; ≈22–30k touched them in the last 30 days

11,072

public repos hold a spec-kit constitution; ≈1,660 touched their specs in 30 days

0

of the tools that write those files is an editor

6

AI routes already live in our repo, wired to a five-provider fallback chain

97 / 53 / 7

findings confirmed / revised / refuted in the round that produced this page

PART I

The decision

1 What we build

GOOD

frontmatter is the editor where you see what the agent changed that you did not ask for, and know what nobody has reviewed yet. Open a spec the agent wrote. Every span changed since a person last reviewed it is tinted; one key reverts any of it; who wrote it and the prompt appear whenever the write carried them. Nothing else in the file moves.

Three corrections to the last version's page one.

- **"Who wrote it" shipped three years ago and did not sell.** iA Writer 7 has dimmed AI-pasted text since 2023-11-30. The Obsidian request to copy it got 65 likes and went quiet within a month; the port has 208 downloads; the best of twelve VS Code extensions has 366 installs; Reddit holds one ask in 107 matched posts, at 2 points.
- **What people upvote is changes they did not ask for and output they cannot review.** Claude Code's diff-review request: 262 reactions; "diff" in 463 of 89,761 issue titles. Requests to *remove* AI authorship marks from commits: 29 and 36. Review what changed. Do not tint who wrote it.
- **"Answer in place" is a small spec-kit feature.** The 62,976 figure was a template artifact: 49,408 of those files are spec-kit's own checklist line *"No [NEEDS CLARIFICATION] markers remain"*. Live markers: ≈1,500 public files, 2.7% of committed specs. superpowers has none. `/speckit.clarify` resolves most in chat.

Why the surface is still ours. Claude Code opens plans as markdown with inline comments (since 2.1.70), VS Code previews range feedback, git-ai records line attribution after the fact, Anthropic's models now carry signed provenance metadata in the EU. Each records a moment. None keeps the state: which spans a person has read, persisting across sessions, teammates and whatever tool wrote the file. That state is a byte-range record beside the file in the user's repo, and keeping it valid while the file changes is what our engine does.

The fifteen-second demo. Run your agent's spec command. Open the spec here: every paragraph tinted, because nobody has read it. Read one; it clears. Revert one; `git diff` shows only that. The panel lists what is still unread — today, next week, or when Amit opens it.

PART II

The market, in one page

2 Who already does what, and where it stops

EXHIBIT 1

Fetched 2026-09-06, denominators attached

	Tool	Real reach	What it does	Where it stops
Spec writers	obra/superpowers	≈75–85k public repos with docs/superpowers/specs or /plans; ≈22–30k active in 30 days; marketplace counter shows 1,009,371 (cached, definition unknown)	Brainstorm → design → writes the spec, commits it, review gate	Terminal. Not an editor. No review state
	github/spec-kit	11,072 public repos with a committed constitution (82% real, not template); ≈1,660 active in 30 days	/speckit.specify · /speckit.clarify (chat-first) · plan · tasks · implement	Writes markdown, then leaves
	Plan mode — Claude Code, Cursor	Out-scores both named tools on Reddit (1,776 / 1,409 / 1,329 points vs ≤38); Cursor's markdown plan editor praised at 167	Plans as markdown files	Where those files live on disk is untested
Review surfaces	Claude Code (VS Code)	24.89M installs	Plans as a full markdown document with inline comments since 2.1.70 (2026-03-06) — it erased the third-party plan-review market (222 installs across three tools)	A moment, per session. Diff-review UI still open at 262 reactions
	git-ai · VS Code 1.118	61,517 installs · default AI co-author trailers	Line-level attribution after the fact; commit trailers	Commit granularity. Nothing persists per span
	SpecKit Companion	9,367 installs	Keeps review state in .spec-context.json per spec — our sidecar pattern, already shipped	Spec-kit only
Editors	Obsidian	Free · Sync \$4 · Publish \$8/site · Commercial \$50/yr	AI arrives as plugins: Copilot 1.83M downloads ("run Claude Code, Codex and OpenCode inside your vault"), Claudian 2.02M	None records what changed or what was reviewed
	Cursor · Zed	~\$20 · free	Strongest agent editing for code; Zed reviews per hunk	Whole-file rewrite; no markdown vault semantics
Browser	Obsidian Web Clipper · Markdown Viewer	1,000,000 · 500,000 Chrome users	Capture and read markdown in the browser	A channel
Hosted sites	mdown.ai · Obsidian Publish	Aimed at non-engineers · \$8/site/mo	Markdown → live website, on their servers	The host role we refuse. We export; they host

The price column, corrected. The editor is worth ₹0. The one proven individual price is sync, ~\$4. The last version called Obsidian's \$50/year commercial licence a second proven price; it *does not provide any functional benefits within the app*, became optional on 2025-02-20, and its buyers are 25+ seat organisations. It says nothing about a five-person team.

3 What was jargon in the original plan, and what survives

EXHIBIT 2

The audit of our own vocabulary, re-run

It said	What it actually was	Verdict
"Context packs", "handover generation"	Free from GitHub and Anthropic now; median-2 launches were the platform baseline anyway	Cut
"Decision-flow renders"	A note-taker's plugin at 0.30× the median, wrong buyer	Cut
"Kickoff prompt", "our CDN files"	A copy-paste step worse than a slash command; a hop that breaks <i>we never hold your documents</i>	Cut
"Self-improving AIOS"	One feedback label from 6,884 decisions	Internal only

It said	What it actually was	Verdict
"Local LLM fallback"	Five <i>cloud</i> free tiers. A constrained small model fabricates rather than refuses (100% in 10 of 13 models)	Defer
"Byte-exactness" as the headline	4 complaints in 12,556	Mechanism, never the pitch
"See what the AI wrote" — authorship	Shipped by iA Writer in 2023. 208 downloads for the port, 366 for the best of twelve extensions, one Reddit ask	Cut as the headline. Keep author when a write supplies it
"496,000 stars" as the channel	Stars are attention to a CLI. Real: low tens of thousands of monthly-active public projects	Resized, above
"Uncopyable by anything that rewrites files"	False at commit granularity — trailers and git-ai do it after the fact	Struck
What changed since you reviewed, kept as a state	The upvoted pain. Nobody persists it	The product — if the two-week test passes
Answer in place	≈1,500 public files, spec-kit only	A feature
File-scoped AI	Table stakes. Two verbs have demand (§4)	Keep, bring-your-own-key, collapsed

PART III

What we ship

4 Keep, cut, defer

EXHIBIT 3

Every proposed component, decided — nine changed since the last version

#	Proposed	Verdict	Why
1	The markdown editor — Live · Edit · Split · Read	KEEP	The free tier. Zed launched with zero AI words; Obsidian's first build had none
2	The design system	KEEP	Replace the sibling project's <code>globals.css</code>
3	Review state — tint since last reviewed, revert, the panel	KEEP — the product	The only thing in Exhibit 1 nobody persists
4	Author and prompt on hover	RESHAPE	Shown when a write carried them: an MCP write, a commit trailer, a git-ai note, signed metadata. At MVP-0 a file the agent wrote with its own tools has no prompt to show — the state is "changed since you reviewed", and the demo must say so
5	Answer [NEEDS CLARIFICATION] in place	KEEP as a feature	≈1,500 live public files. Cheap, spec-kit only
6	The AI writing section, on load	RESHAPE — collapsed by default, no model call on load, a "hide all AI" switch on day one	No incumbent opens a panel at rest; Zed, VS Code and Telegram each shipped a global hide switch after 412-, 30- and 181-reaction issues. Our own record already refused ambient AI; the sketch was a regression
7	AI verbs in the panel	RESHAPE to two	Editor communities request AI at 0.13% of feature titles. Usage data: editing or critiquing text you supplied is 10.6% of all ChatGPT messages and Notion's #1 and #3. Ship "Fix / critique this selection" ; "Summarise this file" provisionally. No translate, tag, outline, explain
8	Comments — drawn in the rail	RESHAPE → repo sidecar, v1.5	Every in-file syntax fails on GitHub: <code>%%</code> and CriticMarkup print verbatim, <code>> [! comment]</code> renders as a plain quote, HTML comments hide but put reviewer bytes inside the artifact. Anchor by byte range + content hash, relocated by the splice locator
9	Bookmarks — drawn in the rail	CUT	Per-user, per-device state with no measured demand
10	Share — drawn in the rail	OPEN	"Permalink + export" needs no state, but a plain folder has no commit and a private repo has no readable link — and our own record's distribution arithmetic assumed a hosted viewer. Decision #3

#	Proposed	Verdict	Why
11	Generate a site or deck	MOVE to MVP-2 as Export	Marp is 5.9% of the top markdown extension's installs; decks are 0.7% of Obsidian downloads; sites-from-a-folder are wanted <i>hosted</i> , which we refuse. Export (HTML, PDF, copy-as-HTML) leads decks 2:1
12	Team tier: who-reviewed-what at \$8	RESHAPE	GitHub Team sells required reviewers and CODEOWNERS on private repos at \$4. What small teams are gated on elsewhere: seats, history retention, shared-vault sync. Reprice \$8
13	Generate all files · kickoff prompt · CDN · local model	CUT / DEFER	Unchanged from the last version

5 The three tiers

EXHIBIT 4

MVP-0 pilot · MVP-1 pro · MVP-2 max

	MVP-0 · Pilot — 10 weeks, free forever	MVP-1 · Pro — teams	MVP-2 · Max
Engine	NF-1 + NF-3 fixed · CI with one deliberate red run · engine wired behind every write · real byte budget	Vault-wide refactor with a reviewable diff, ambiguity refused	Sync, provably safe — git merge, never a CRDT
Editor	Four modes · tree with the unreviewed bar · tabs · quick-switch · palette · search · properties · tags	Conflict view	Multi-device
Review state	Span-level "changed since reviewed" by content hash, in a sidecar in <i>their</i> repo · tint · one-key revert · review panel · author when a trailer, note or MCP write supplies it	Shared review state across a repo — what each person has read · comments as a sidecar (v1.5)	Version-history retention
Answer in place	[NEEDS CLARIFICATION] markers, inline	—	—
AI	Collapsed panel, bring-your-own-key only , two verbs · hide-all switch · idea mode as it ships · structured output in the port	MCP server so the agent's writes carry the prompt · hosted AI metered, hard-capped, never default	Document gates in CI
Output	—	—	Single-document Export: HTML · PDF · copy-as-HTML. Decks if a pilot user asks
Channels	Free Obsidian plugin, week 2 — kill signal under 200 installs in 14 days	Chrome extension: open any .md URL here	—
Absent, on purpose	Sign-up · billing · sync · mobile · hosting · hosted AI · local model · generation	Mobile · plugins · a chat sidebar	Plugins, still
Exit criterion	6 of 10 strangers keep it after two weeks — observed, not felt (\$10)	A team pays within 20 qualified conversations	Sync passes the corpus byte-identical

Before any stranger touches it, three things that do not exist today must: bring-your-own-key (settings expose one boolean; every request bills us), structured output in the LLM port (it returns a bare string), and CI (four gates have reported green while blind). **And one decision the round exposed:** the byte path for the web trial — whether "connect a repo" fetches document bytes client-side with the user's token or through our server. Nobody has written it down, and "zero document bytes in our control plane" depends on the answer.

6 How it is built — the plugin question, answered

The cheapest shape is the one where **we run no model, hold no document, and sit inside the tools the user already pays for.**

EXHIBIT 5

The delivery surfaces

Surface	Role	Cost to us	When
Web app	The trial. Visit a URL, connect a repo. Byte path to be decided (\$5)	~\$0	MVP-0
Desktop (Tauri)	The daily surface. Local folders; offline everything but language AI. Same build as the web app	\$99/yr Apple; Windows signing \$150–400/yr with a hardware token	MVP-0
Review state from content hashes	Works for files written by <i>any</i> tool — the toolchains write with the agent's own tools, not ours	~\$0	MVP-0

Surface	Role	Cost to us	When
Their agent, via MCP	The write carries the prompt. An assumption: agents own Write/Edit and the toolchains instruct direct writes; whether one routes through us is untested	~\$0	MVP-1
Bring-your-own-key panel	Their key, their machine → provider. Never through us	~\$0	MVP-0
Obsidian plugin	Distribution only	4 days	Week 2
Chrome extension	Distribution. 1.5M people already read markdown in the browser	~1 week	MVP-1
Vendor CDN · hosting	—	—	Never

So, "do we ask them to install a plugin": no. The plugin is a channel. The product is the web app and the desktop build; the AI is theirs.

The rules that keep it cheap and honest: documents never enter our control plane · nothing runs in the background, so spend follows use · **cap output, not input** — output is 75.8% of spend · the byte↔UTF-16 seam is the highest-risk thing we own · review state and comments anchor to **bytes plus a content hash and are relocated by the splice locator, never by our journal** — git, other editors and agents bypass the journal, and the shipped competitor's .spec-context.json is the pattern we must be visibly better than.

PART IV


The screens

7 As they will actually look

Built from the two hand-drawn layouts, with real content. Clickable prototype: [claude.ai/code/artifact/f6baec30-9d0f-4e4f-a231-50f178f0e65e](#) — also in the repo at docs/prototype/frontmatter-prototype.html.


EXHIBIT 6

SO · The launcher — blank, four templates, existing edits sorted unreviewed-first



The bar on every card is the product showing itself before anything is open. Templates are markdown skeletons: spec, decision record, handover, changelog.

S2 · The editor — tree, coloured tabs, four modes, the review tint, and the AI panel collapsed



Changed-since-reviewed text carries a light tint and nothing else. Click a span: **when it changed, which bytes, the author from the commit trailer — and, honestly, "prompt not recorded" for a file the agent wrote outside frontmatter**. The AI panel is a collapsed line: nothing runs until you ask, ⌘J opens it, Settings hides it entirely.

EXHIBIT 8

S3 · Answer in place — a spec-kit feature, kept because it is cheap

fm

spec.mdplan.mdtasks.md

ShareSearchsm

TREE

acme-billing

specs

001-invoicing

spec.mdplan.mdtasks.md

.specify

AGENTS.md

LiveEditSplitRead

written by spec-kit via Claude Code · 11 min ago · 3 questions for you

feature: 001-invoicing

status: draft

generated-by: speckit.specify · 2026-09-06

Feature Specification: Invoicing

Functional requirements

- FR-001: System MUST let an admin create an invoice from a customer record with line items, tax and a due date.
- FR-002: System MUST render the invoice as a PDF that is byte-identical across renders for the same input.
- FR-006: System MUST authenticate users via

A SSO through the GitHub App – no second identity path
- FR-007: System MUST retain invoices for

NEEDS CLARIFICATION 7 years – Indian Companies , Answer
- FR-009: System MUST charge Indian cards

NEEDS CLARIFICATION once, or retry on failure?

Non-goals

Multi-currency invoicing, refunds automation, and any accounting export are out of scope for this feature.

Answering writes → A: _ into the file at that byte range, exactly where spec-kit reads it back. 1 line changed, 0 lines moved. The requirements the agent wrote stay tinted until you have read them.

7 unreviewed spans · 92%2 questions open

git · +1 -1 · spec.mdLF · UTF-8

Questions in this spec2

FR-007 · retention periodopen

FR-009 · card retry policyopen

FR-006 · auth methodanswered

Document outline

Functional requirements5 · 5unreviewed

Non-goals1 · 1

Document history

nowyou · answered FR-006

11mspec-kit · /speckit.specify · 2,140 bytes

Mark all readShortcuts

AI edit · #J

EXHIBIT 9

S4 · The review panel — catching up on a file you did not watch being written

fm

0004-sync.mdauth.mdpricing.md

ShareSearchsm

TREE

acme-billing

specs

adr

0003-scope.md0004-sync.mdn

notes

pricing.md

LiveEditSplitRead

acme-billing / adr / 0004-sync.md

title: ADR-0004 Sync without a CRDT

status: proposed

date: 2026-09-02

ADR-0004 · Sync without a CRDT

Context

Two devices edit the same markdown file. Most sync products reach for a CRDT because convergence is guaranteed. Convergence is not the property we want: CRDTs interleave concurrent edits character by character, so two people finishing the same sentence produce a merged sentence neither wrote.

Decision

Git three-way merge, with divergence surfaced as a reviewable conflict. On the same inputs it preserves Windows line endings and a missing final newline, and it stops exactly where a careful writer would want to stop.

Consequences

- Conflicts happen, and users hate conflicts. We publish the conflict rate as a shipped metric with a budget.
- We never guess which version was meant. That is the whole point.

3 unreviewed spans · 88%231 words

git · clean · mainCRLF · UTF-8

3 UNREVIEWED IN THIS FILE

Context · 312 bytes · sonnet-5 · 2h

Two devices edit the same markdown file. Most sync products reach for a CRDT because...

a · r · s · revert · skip

Decision · 218 bytes · sonnet-5 · 2h

Git three-way merge, with divergence surfaced as a reviewable conflict...

Consequences · 96 bytes · sonnet-5 · 2h

Conflicts happen, and users hate conflicts. We publish the conflict rate...

Provenance is a state, not a moment. These stay marked until a person accepts or reverts them — today, next week, or when a teammate opens the file.

Every changed span in order, j k to move, a r s to accept, revert or skip, a counter that reaches zero. **This is the state nobody else keeps.**

EXHIBIT 10

S2 · Split — source and rendered, the same bytes

fm

pricing.md xauth.md x

ShareSearchsm

< TREEFP+

acme-billing

- notes
- pricing.md
- specs
- adr

LiveEditSplitRead

acme-billing / notes / pricing.md · source = rendered, same bytes

title: Pricing

status: decided

tags: [pricing, obsidian, teams]

Pricing

The editor is free. Obsidian settled that: the whole editor at \$0, sync at \$4, a \$50/user/year commercial licence. Everyone priced around it.

What we charge for

- **Team** - \$8/user/month. Provenance across a shared repo, review state across people, admin.

- **Sync** - \$4/user/month, later, when it is provably safe.

The arithmetic

At \$8/seat, ****114 paying seats**** covers the monthly nut of ₹1.09L - about twenty small teams. Churn on a team licence is far lower than on a personal one.

title: Pricing

status: decided

tags: [pricing, obsidian, teams]

Pricing

The editor is free. Obsidian settled that: the whole editor at \$0, sync at \$4, a \$50/user/year commercial licence. Everyone priced around it.

What we charge for

- **Team** — \$8/user/month. Provenance across a shared repo, review state across people, admin.
- **Sync** — \$4/user/month, later, when it is provably safe.

The arithmetic

At \$8/seat, **114 paying seats** covers the monthly nut of ₹1.09L — about twenty small teams. Churn on a team licence is far lower than on a personal one.

2 unreviewed spans · 31%96 words

git · clean · mainLF · UTF-8

EXHIBIT 11

S11 · Settings — the hide-all-AI switch, and the key field that does not exist yet

fm

Settings

sm

Settings

APPEARANCE

Theme

System

Editor typeface

Google Sans Code

KEYMAP

Default

AI PROVIDER

☐

Hide all AI features. Removes every AI control from the interface. The review tint stays — it is file state, not AI

Provider

Anthropic

Your API key

.....

Stored in your OS keychain. Requests go from this machine to the provider. They do not pass through us.

Model for language work

claude-sonnet-5

Structural work — links, slugs, frontmatter fills — never calls a model.

☐

Use frontmatter's hosted AI instead (metered, hard-capped, never the default)

PROVENANCE

☒

Mark machine-written text until a person reviews it

☒

Show the unreviewed bar in the file tree

Record

byte rangepromptmodeltime

session id

Written to .frontmatter/provenance.jsonl beside your documents, in your repo. Plain text. Readable without us.

EXHIBIT 12

S8 · Refactor preview (MVP-1) — the engine made visible

fm

Rename tag #auth → #authentication

Cancel

Apply 6 files

8 FILES · 6 CHANGE · 2 REFUSED

specs/auth.md2 hunks

specs/billing.md1

specs/api.md1

adr/0003-scope.md1

notes/pricing.md1

README.md1

notes/2026-08-12.mdrefused

archive/old-auth.mdrefused

specs/auth.md · frontmatter · bytes 41–451 of 2

4 tags: [auth, github-app, security]

4 tags: [authentication, github-app, security]

acceptskip

accept all in file

specs/auth.md · body · bytes 1,982–1,9872 of 2

38 See also [[billing]] and the

39 #auth notes from the kickoff.

39 #authentication notes from the kickoff.

acceptskip

REFUSED · 2 FILES

notes/2026-08-12.md — #auth appears inside a fenced code block at line 17; rewriting code is a guess we will not make. archive/old-auth.md — the tag sits in flow form tags: auth with no brackets; the target is ambiguous. Both files are untouched.

Every file that will change is listed before a byte moves. Accept all, per file, or per hunk. Ambiguity is refused, never guessed — and the refusal says why.

EXHIBIT 13

S12 · Export (MVP-2) — a site or page from the folder, never hosted by us

fmGenerate · acme-billing ×

OUTPUT

Static site · tree as navigation

One shareable page

Slide deck · H2 = slide

WHAT GOES IN

specs/ · adrl · README.md

notes/ excluded

☐ Skip documents with unreviewed machine text

THEME

Plain · one column

Docs · sidebar

Anything custom is site, css in your repo.

DEPLOY

Vercel · your account

GitHub Pages

Netlify

We never host it. Files land in . /site in your repo. Deploy runs on your account, under your terms.

acme-billing.vercel.app/specs/auth

Specs

Authentication

Billing

API

Decisions

0003 Scope

0004 Sync

Authentication

We use the GitHub App installation flow rather than an OAuth app, because the App grants permission per repository instead of across the whole account.

The installation token is scoped to the repositories the user selected and expires after one hour, which means a leaked token has a bounded blast radius.

What we never do. Ask for the repo scope. Keep a second identity path.

generated by frontmatter · every paragraph reviewed by a person · source: specs/auth.md

A site is another projection of the same files. No site editor, no theme builder — edit the markdown, regenerate.

PART V

Money, market, and the plan

8 Price, with the churn arithmetic attached

EXHIBIT 14

What we charge, corrected against what teams are gated on

Tier	Price	What it is
Individual	Free, forever, no limits	The editor, review state, revert, answer-in-place, BYO key. This is the distribution
Pro · Team	\$4–5/user/month — priced against GitHub Team \$4, Obsidian Sync \$4, HackMD \$5, Confluence \$6.70	Shared-vault sync, version-history retention, shared review state, comments, MCP. Not "who reviewed what" — GitHub sells that at \$4
Max	Later, per team	Document gates in CI, retention beyond a year, admin
Hosted AI	Metered, at cost plus a margin	Optional, never the default

The arithmetic. Monthly nut ~₹1.09L, covered by 54–78 consulting hours. Infrastructure ~\$0 at 100 users, ~\$390/month at 10,000; support is the real cost, ~90 founder-hours a month at scale. **"114 seats at \$8" is withdrawn:** the feature it priced is GitHub's at \$4, and accounts under \$50 a month sit in the worst-retaining band there is — top-quartile annual gross retention 60–70% across 2,100 businesses, 23% for AI-native products. Holding 114 seats means replacing 34–46 a year at best. Nobody comparable publishes retention; Obsidian says on record *"we don't know what suddenly causes someone to churn."* **The tier is a hypothesis until the 20 conversations in §11 return seat counts and what teams are gated on.**

TEST

Give away what the category has proven is free. Charge for what small teams are gated on elsewhere — seats, history, sync — and never for a feature GitHub bundles. Downgrading loses features, never files.

9 What grabs users, and how we reach them

The message is now three arms, tested in weeks 1–2, not one line chosen by us: *See what the AI wrote* · **See what the agent changed that you did not ask for** · **Know what nobody has reviewed yet**. Reddit's top twenty "AI wrote" posts contain zero provenance asks; the upvoted pain is the second and third. The proof line under any of them: *git diff after an edit shows your change and nothing else. Tested on 8,513 real files, every release.*

Lead with the refusals — this audience is sceptical of AI tooling: no plugins, no code execution, no holding your files, no lock-in, **and a switch that hides every AI feature**, which the incumbents each shipped only after a backlash. Never *revolutionary, seamless, AI-native, or byte-preserving* in a headline.

EXHIBIT 15
Channels, resized to real usage, in the order we use them

When	Channel	Size, measured	What we do
Week 2	Free Obsidian plugin — the live-preview fix	501 likes, 117 posts, still open since 2022. Counter-evidence: Obsidian patches live preview monthly; the nearest plugin-level render fix has 1,454 downloads	Ship it; kill signal under 200 installs in 14 days
Week 2 →	superpowers first, spec-kit second, plan-mode files third	≈22–30k · ≈1,660 active public projects a month · plan-mode size untested	A README line, an MCP entry, "open the spec here" in their docs. Maintainer PRs are outbound acts — we ask before sending
Now →	Our own writing	—	The research behind this document. Nobody has published it
MVP-0 exit	Show HN	Zed's scored 43, Cursor's 9	After 6 of 10 strangers keep it. One shot
MVP-1	Chrome extension · Product Hunt	1.5M browser readers	Backlinks and a spike
MVP-2	Exported pages with an honest mark	Not a channel until a deploy count exists	Withdrawn from "always"

How we use the users we get. The ten pilot strangers are the interviews we have never done. Every accept and revert is a preference signal — logged **only** in a local sidecar they can read, because the same architecture that keeps their documents out of our hands keeps their usage out too (§10). The plugin's installers are the list we do not have.

10 Challenges, and what we do about each

EXHIBIT 16
Ranked by what would stop us — two rows new this round

Challenge	What we do
Review state turns out to be a nice-to-have	Two-week test, three headline arms, kill gate: fewer than 4 of 10 call it useful unprompted
Anthropic ships the review surface — plan-as-markdown with comments already exists since 2.1.70; the diff-review UI request is open at 262 reactions	Persistence is what they do not build: across sessions, across people, across tools. If they build that too, we are a feature, and we should know by week 12
The MVP-0 exit gate is unobservable — under our own architecture we cannot see who "keeps it", and Obsidian cannot either	A written measurement decision before week 10: opt-in, local sidecar counts they can read; never a vendor usage database
Nobody pays for the team tier	Run the 20 conversations in weeks 1–2, not at MVP-1. Ask what they are gated on, and how many seats they are
Zed or Cursor add markdown vault semantics — one sprint	Depth in markdown, and the state they do not keep
Distribution	Three cheap channels, resized to real numbers. The private-repo multiplier is unmeasured — ask the ten strangers
The 83% refusal is a class, not a bug	Falsified if a patched corpus still refuses over 10 of 7,969, by day 24
Every AI call bills us	BYO key before any stranger
Client work crowds out product	Capped at 78 hours a month, written down
GST	Reverse charge has no turnover floor

11 Ninety days, and the decisions

EXHIBIT 17

An observable outcome every fortnight

Weeks	What happens	Outcome
0	Rotate the two access tokens · decide the buyer · write the web byte-path and the measurement decision	Done, written down
1–2	Four tests, build nothing : 10 developers shown the review mock with both framings (who-wrote-it vs changed-since-reviewed) and asked which toolchain they use · plugin shipped · 5 people who bill for documents asked about approval and sign-off , not provenance · 3 landing pages, one per headline arm · 20 team conversations on seats and gating	Kill signals at day 14
3	Go / no-go	A written decision either way
3–4	CI with a deliberate red run · NF-3, then NF-1 · prototype content-hash review state on one spec-kit and one superpowers repo	CI fails a broken commit. Refusals fall from 83%. The tint works on files we did not write
5–7	Review state store and rendering · BYO key · structured output	A stranger's key, a stranger's repo, tinted spans
8–9	Revert · review panel · answer-in-place · engine wired behind every write	The fifteen-second demo, as written in §1
10–12	Table stakes · ten strangers, their own repos and agents	6 of 10 keep it, by the measure written in week 0

EXHIBIT 18

The decisions

#	Decision	By
1	Accept the reshape — review state as the product, authorship as a detail, answer-in-place as a feature, the team tier as a hypothesis	The meeting
2	Which buyer first — the developer (reachable; AI-native products under \$50 retain 23% of revenue a year) or the person liable when a document is wrong	Week 1
3	Share — a hosted read-only viewer (our old record's assumption, and a hosting obligation) or permalink-plus-export (stateless, but nothing for a plain folder or a private repo)	Week 2
4	The two unrotated access tokens — an action, not a decision	Today

RISK

What would make me say stop. Fewer than 4 of 10 developers calling review state useful under either framing, and the 20 team conversations returning nothing they are gated on that GitHub does not already sell. Then there is no product here — only an engine, a services business, and a free plugin people like. That is not a failure; it is a smaller, truer version of the same work, and it pays sooner.

PART VI

The rest, in one page

12 The engine — what exists, what is broken, what must be optimised

EXHIBIT 19

Built · broken · optimised

Built	25,407 lines, 1,575 tests. 8,513 real files from strangers' vaults, byte-pinned, zero corruption. <code>git diff</code> after an edit shows only the edit — tested every release. Six AI routes, five-provider chain, idea mode
Broken	NF-1 — a column-zero list item in frontmatter refuses 83% of real vaults, 4 days · NF-3 — a bare-CR fence adds a second frontmatter block, 3 days · no CI , and four gates have reported green while blind, 1 day · engine unwired — one symbol from one of thirteen files reaches product code, 8 days · byte budget is an echo , 2 days
Optimise	Incremental parsing — vendored <code>@lezer/markdown</code> , so review spans survive every keystroke cheaply · the offset map — bytes to UTF-16, the highest-risk seam; crossing it unmapped corrupts Chinese, Japanese, Korean and Arabic silently · search — today the whole vault ships to the client, 77 MB parsed per cold start; move to server full-text, client only for the open file · refusal and conflict rates published as metrics with budgets

Published measures

Time to first keystroke, cold · typing latency on a 10,000-word document · corpus refusal rate (83% → near zero) · merge conflict rate

13

What we render, and the markdown rules

EXHIBIT 20

Custom rendering, and what we never do to a file

Modes	Live · Edit · Split · Read. Nested constructs inside list items render correctly — the 501-like bug, shipped as the free plugin first
We render	Callouts > [!kind] · Mermaid, rendered not extended · MADR and Nygard decision records · RFC 7322 · Keep a Changelog 1.1.0 · frontmatter as a properties panel
We do not pretend	Runbook and PRD have no standard body. We offer a structure and say so — every "PRD standard" claim on the market is unsourced
The rule	We add nothing to markdown that breaks it somewhere else. A file we touch renders on GitHub, in Obsidian, in a plain editor, next year
How	Prose annotations as callouts — no closing marker to lose; an unclosed fence swallows the document · tags and links written back in the shape the file already uses , never converted · wikilinks and standard links both read · review state and comments in a plain-text sidecar beside the file, readable without us
What we join	Read AGENTS.md and CLAUDE.md, never compete with them · be an MCP server the user's agent connects to · we are a git client, never our own versioning · Obsidian conventions preserved

14

Where we stand

EXHIBIT 21

Strengths · weaknesses · opportunities · threats

Strengths	The engine is real and correct where the incumbents admit they are not · 88.6% of a shipping editor exists · six AI routes live · two founders funded by services, not a clock
Weaknesses	Zero users, zero interviews · no CI, no BYO key, no structured output · 83% of vaults refused today · every channel is borrowed · the paid tier has no evidence yet · we have published three numbers that failed a second check
Opportunities	Tens of thousands of active projects generating specs with nowhere to review them · review <i>state</i> is the one thing no incumbent persists · a 501-like bug we can fix in a week · a hide-all-AI switch the incumbents shipped only after backlash
Threats	Anthropic ships persistent review — plan-as-markdown exists since 2.1.70, diff review is open at 262 reactions · Zed or Cursor add vault semantics in a sprint · Obsidian ships first-party AI editing · attention runs out before revenue

15

Operating rules

EXHIBIT 22

The lines we hold

D2C / B2B	Free to build the audience; paid to build revenue. No enterprise feature until a customer refuses to pay without it — SSO at ~20 seats when asked, SOC 2 above ~50 or for a regulated buyer, a merchant of record for EU/US procurement
How users hear from us	We collect an email only when an obligation is created — a payment, an invite, a recovery request — never at first run. Breaking changes: two weeks' notice in product. Price changes: email before the next charge. Every notice lives in a quiet in-product ledger, so "we told you" is verifiable by them
How they reach us	One address, a published response window, and the free tier has no service commitment , said on the page
What we refuse to measure	Session replay, keystroke telemetry, document content — the DOM would be their private document. Usage counts live in a local sidecar they can read
Mobile	~15% of complaints are mobile. Not in v1. When it comes, a reader and reviewer, never a byte-exact editor on a phone keyboard
Accessibility · non-English	WCAG 2.2 AA, all-or-nothing — our body-faint token fails at 1.984:1 and is fixed in MVP-0. The review tint is never colour alone: the panel and the keyboard carry the same information. CJK: the byte map, a bigram search tokeniser, and a written decision on v1 scope
Security	No plugins, no code execution — the largest single mitigation we have. The AI never gets ambient repo access. Refuse a malformed repo rather than repair it. Reverse charge under CGST §24(iii) has no turnover floor

16 What I think

- **The engine is the best thing either of us has built.** Two verification rounds did not touch it.
- **We were wrong about the pitch three times** — bytes, then authorship, then a generator — and each time the evidence was public and unopened. What survives is narrower: a state nobody keeps, on files everyone now generates, in a surface nobody owns. It is a product only if ten strangers say so in two weeks.
- **The open question is no longer distribution. It is whether anyone pays.** A free product with a channel and no evidence for the paid tier. The twenty conversations settle it, in week one.