

The pilot plan: what we build first, what we cut, and how we sell it

Amit, this is the plan for our meeting. I checked it against 157 sources this week and had to correct 60 things, three of them on this very page. So the numbers here were opened, not remembered. Where I was wrong before, I have said so in the same place.

KEY FIGURES

75–85k

public repos hold superpowers specs and plans. 22–30k of them were touched in the last 30 days

11,072

public repos hold a spec-kit constitution. About 1,660 touched their specs in 30 days

0

of the tools that write those files is an editor

6

AI routes already live in our repo, on a five-provider fallback chain

97 / 53 / 7

findings confirmed, revised, refuted in the round behind this page

PART I

The decision

1 What we build

GOOD

frontmatter is the editor where you see what the agent changed without being asked, and know what nobody has reviewed yet. Open a spec the agent wrote. Everything changed since a person last read it is tinted. One key reverts any of it. Who wrote it, and the prompt, show up whenever the write carried them. Nothing else in the file moves.

Three things I got wrong on this page last time.

- **"See what the AI wrote" was shipped three years ago, and nobody bought it.** iA Writer 7 has dimmed AI-pasted text since November 2023. The Obsidian request to copy it got 65 likes and went quiet within a month. The port has 208 downloads. The best of twelve VS Code extensions doing the same thing has 366 installs. On Reddit there is one ask for it in 107 matching posts, with 2 points.
- **What people actually upvote is changes they did not ask for, and output they cannot review.** In Claude Code's own tracker the diff-review request has 262 reactions and "diff" is in 463 of 89,761 issue titles. Requests to *remove* AI authorship marks from commits got 29 and 36. So: show what changed. Do not tint who wrote it.
- **"Answer in place" is a small spec-kit feature, not the product.** The 62,976 figure I quoted was a template artefact. 49,408 of those files are spec-kit's own checklist line, *"No [NEEDS CLARIFICATION] markers remain."* Real open markers: about 1,500 public files, 2.7% of committed specs. superpowers has no such marker, and /speckit.clarify settles most of them in chat anyway.

Why the reading surface is still ours. Claude Code now opens plans as markdown with inline comments. VS Code previews range feedback on markdown. git-ai records who wrote which line, after the fact. Anthropic's models carry signed provenance metadata in the EU. Each of these records a moment. None of them keeps the state: which spans a person has read, kept across sessions, across teammates, across whatever tool wrote the file. That state is a byte-range record sitting next to the file in the user's own repo, and keeping it valid while the file changes is exactly the job our engine does.

The demo, fifteen seconds. Run your agent's spec command. Open the spec here. Every paragraph is tinted, because nobody has read it. Read one, it clears. Revert one, `git diff` shows only that. The panel on the right lists what is still unread, today, next week, or when you open it.

PART II

The market

2 Who already does what, and where each one stops

EXHIBIT 1

Fetched 2026-09-06, with denominators

	Tool	Real reach	What it does	Where it stops
Spec writers	obra/superpowers	75–85k public repos with docs/superpowers/specs or /plans, 22–30k active in 30 days. The marketplace counter says 1,009,371, cached, definition unknown	Brainstorm, design, write the spec, commit it, gate it for review	A terminal tool. Not an editor. No review state
	github/spec-kit	11,072 public repos with a committed constitution, 82% real. About 1,660 active in 30 days	/speckit.specify, /speckit.clarify (asks in chat), plan, tasks, implement	Writes markdown, then leaves
	Plan mode in Claude Code, Cursor	Out-scores both named tools on Reddit (1,776 / 1,409 / 1,329 points against 38 or fewer). Cursor's markdown plan editor is praised at 167	Plans as markdown files	Where those files live on disk is untested
Review surfaces	Claude Code, VS Code extension	24.89M installs	Plans as a full markdown document with inline comments since 2.1.70, March 2026. This wiped out the third-party plan-review market: 222 installs across three tools	A moment, per session. The diff-review request is still open at 262 reactions
	git-ai, VS Code 1.118	61,517 installs; AI co-author trailers on by default	Line-level attribution after the fact, commit trailers	Commit granularity. Nothing persists per span
	SpecKit Companion	9,367 installs	Keeps review state in a .spec-context.json per spec. Our sidecar idea, already shipped	Spec-kit only
Editors	Obsidian	Free. Sync \$4. Publish \$8 per site. Commercial \$50 a year	AI comes as plugins: Copilot 1.83M downloads ("run Claude Code, Codex and OpenCode inside your vault"), Claudian 2.02M	None records what changed or what was reviewed
	Cursor, Zed	About \$20, free	The strongest agent editing for code. Zed reviews per hunk	Whole-file rewrite. No markdown vault semantics
Browser	Obsidian Web Clipper, Markdown Viewer	1,000,000 and 500,000 Chrome users	Capture and read markdown in the browser	A channel for us
Hosted sites	mdown.ai, Obsidian Publish	Aimed at non-engineers; \$8 a site a month	Markdown to a live website, on their servers	The host role we refuse. We export, they host

On price, corrected. The editor is worth ₹0 in this category. The one thing individuals have proven they pay for is sync, about \$4. Last time I called Obsidian's \$50 commercial licence a second proven price. It is not. It *does not provide any functional benefits within the app*, became optional in February 2025, and its buyers are organisations with 25 or more seats. It says nothing about a five-person team.

3 What was just AI vocabulary, and what survives

EXHIBIT 2

Our own words, checked

We said	What it turned out to be	Verdict
"Context packs", "handover generation"	Free now from GitHub and Anthropic. And the "median 2 points" graveyard was the Hacker News baseline all along	Cut
"Decision-flow renders"	A note-taker's plugin at 0.30 of the median, and the wrong buyer	Cut
"Kickoff prompt", "our CDN files"	A copy-paste step worse than a slash command, and a hop that breaks "we never hold your documents"	Cut
"Self-improving AIOS"	One feedback label from 6,884 decisions	Internal only

We said	What it turned out to be	Verdict
"Local LLM fallback"	Five cloud free tiers, not local. A constrained small model makes things up rather than refusing, in 10 of 13 models	Defer
"Byte-exactness" as the headline	4 complaints in 12,556	Mechanism, never the pitch
"See what the AI wrote"	Shipped by iA Writer in 2023. 208 downloads for the port, 366 for the best extension, one Reddit ask	Cut as the headline. Show the author when a write supplies it
"496,000 stars" as our channel	Stars are attention to a CLI. Real usage is low tens of thousands of active projects a month	Resized, above
"Uncopyable by anything that rewrites files"	False at commit level. Trailers and git-ai do it after the fact	Struck
What changed since you reviewed, kept as a state	The pain people upvote. Nobody persists it	The product, if the two-week test passes
Answer in place	About 1,500 public files, spec-kit only	A feature
File-scoped AI	Table stakes. Two verbs have real demand	Keep, your own key, collapsed

PART III

What we ship

4 Keep, cut, defer

EXHIBIT 3

Every proposed piece, decided. Nine changed since last time

#	Proposed	Verdict	Why
1	The editor: Live, Edit, Split, Read	Keep	The free tier. Zed launched with zero AI words on its page. Obsidian's first build had none
2	The design system	Keep	Replace the sibling project's <code>globals.css</code>
3	Review state: tint since last read, revert, the panel	Keep. This is the product	The only thing in Exhibit 1 nobody persists
4	Author and prompt on hover	Reshape	Shown when the write carried them: an MCP write, a commit trailer, a git-ai note, signed metadata. At MVP-0 a file the agent wrote with its own tools has no prompt to show. The state is "changed since you reviewed", and the demo has to say so
5	Answer [NEEDS CLARIFICATION] in place	Keep as a feature	About 1,500 live public files. Cheap. Spec-kit only
6	The AI writing section, open on load	Reshape: collapsed, no model call on load, a "hide all AI" switch from day one	No incumbent opens a panel at rest. Zed, VS Code and Telegram each shipped a hide switch after 412, 30 and 181-reaction issues. Our own record had already refused ambient AI; the sketch was a step back
7	AI verbs in the panel	Reshape to two	Editor communities ask for AI in 0.13% of feature titles. Usage data says editing or critiquing text you supplied is 10.6% of all ChatGPT messages, and Notion's #1 and #3. Ship "fix or critique this selection" . "Summarise this file" provisionally. No translate, tag, outline or explain
8	Comments, drawn in the rail	Reshape: a sidecar in the repo, v1.5	Every in-file syntax fails on GitHub. <code>%%</code> and CriticMarkup print as literal text, <code>> [! comment]</code> renders as a plain quote, HTML comments hide but put reviewer bytes inside the artefact. Anchor by byte range plus content hash, moved by the splice locator
9	Bookmarks, drawn in the rail	Cut	Per-user, per-device state with no measured demand
10	Share, drawn in the rail	Open	"Permalink plus export" needs no state, but a plain folder has no commit and a private repo has no readable link. And our old record's distribution maths assumed a hosted viewer. Decision 3
11	Generate a site or a deck	Move to MVP-2, as export	Marp is 5.9% of the top markdown extension's installs. Decks are 0.7% of Obsidian downloads. Sites from a folder are wanted hosted, which we refuse. Export (HTML, PDF, copy as HTML) beats decks 2 to 1

#	Proposed	Verdict	Why
12	Team tier: who reviewed what, at \$8	Reshape	GitHub Team sells required reviewers and code owners on private repos at \$4. What small teams are actually gated on elsewhere: seats, history retention, shared-vault sync. Repriced in §8
13	Generate all files, kickoff prompt, CDN, local model	Cut or defer	Unchanged

5 The three tiers

EXHIBIT 4

MVP-0 pilot, MVP-1 pro, MVP-2 max

	MVP-0, the pilot. 10 weeks, free forever	MVP-1, pro. Teams	MVP-2, max
Engine	NF-1 and NF-3 fixed. CI, with one deliberate red run. Engine wired behind every write. A real byte budget	Vault-wide refactor with a reviewable diff, ambiguity refused	Sync, provably safe. Git merge, never a CRDT
Editor	Four modes. Tree with the unreviewed bar. Tabs, quick-switch, palette, search, properties, tags	Conflict view	Multi-device
Review state	Span-level "changed since reviewed" by content hash, in a sidecar in <i>their</i> repo. Tint. One-key revert. The review panel. Author when a trailer, note or MCP write supplies it	Shared review state across a repo: what each person has read. Comments as a sidecar (v1.5)	Version-history retention
Answer in place	[NEEDS CLARIFICATION] markers, inline		
AI	Collapsed panel, your own key only , two verbs. Hide-all switch. Idea mode as it ships today. Structured output in the port	MCP server , so the agent's writes carry the prompt. Hosted AI metered and hard-capped, never the default	Document gates in CI
Output			Single-document export: HTML, PDF, copy as HTML. Decks if a pilot user asks
Channels	Free Obsidian plugin, week 2. Kill signal: under 200 installs in 14 days	Chrome extension: open any .md URL here	
Absent, on purpose	Sign-up, billing, sync, mobile, hosting, hosted AI, local model, generation	Mobile, plugins, a chat sidebar	Plugins, still
Exit test	6 of 10 strangers keep it after two weeks, measured, not felt (§10)	A team pays within 20 real conversations	Sync passes the corpus byte-identical

Before any stranger touches it, three things that do not exist today have to. Bring-your-own-key (settings expose one boolean today, and every request bills us). Structured output in the LLM port (it returns a bare string). And CI (four of our gates have reported green while blind). Also one thing the research exposed: nobody has written down the byte path for the web trial. Does "connect a repo" fetch document bytes in the browser with the user's token, or through our server? "Zero document bytes in our control plane" depends on the answer.

6 How it is built, and the plugin question

The cheapest shape is the one where we run no model, hold no document, and sit inside the tools the user already pays for.

EXHIBIT 5

The delivery surfaces

Surface	Role	Cost to us	When
Web app	The trial. Visit a URL, connect a repo. Byte path to be decided (§5)	About \$0	MVP-0
Desktop, Tauri	The daily surface. Local folders. Works offline for everything except language AI. Same build as the web app	\$99 a year for Apple notarisation. Windows signing \$150–400 a year, now with a hardware token	MVP-0
Review state from content hashes	Works for files written by any tool, which matters because the toolchains write with the agent's own tools, not ours	About \$0	MVP-0
Their agent, over MCP	The write carries the prompt. This is an assumption: agents own Write and Edit, and the toolchains tell them to write files directly. Whether any agent routes through us is untested	About \$0	MVP-1

Surface	Role	Cost to us	When
Bring-your-own-key panel	Their key, their machine to their provider. Never through us	About \$0	MVP-0
Obsidian plugin	Distribution only	4 days	Week 2
Chrome extension	Distribution. 1.5M people already read markdown in a browser	About a week	MVP-1
Vendor CDN, hosting			Never

So, do we ask people to install a plugin? No. The plugin is a channel. The product is the web app and the desktop build, and the AI is theirs.

A few rules keep this cheap and honest. Documents never enter our control plane. Nothing runs in the background, so spend follows use, not time. Cap output, not input: output is 75.8% of spend. The byte-to-UTF-16 seam is the riskiest thing we own. And review state and comments anchor to bytes plus a content hash and get moved by the splice locator, never by our own journal, because git, other editors and agents all bypass the journal. SpecKit Companion's `.spec-context.json` is the pattern we have to be visibly better than.

PART IV

The screens

7 As they will actually look

Built from the two hand-drawn layouts, with real content. Clickable prototype: [claude.ai/code/artifact/f6baec30-9d0f-4e4f-a231-50f178f0e65e](#), also in the repo at docs/prototype/frontmatter-prototype.html.

EXHIBIT 6

S0, the launcher. Blank, four templates, existing edits sorted unreviewed-first

fm

Search every document, every project

Open folder

New

START SOMETHING

Blank document

empty file · yours from byte one

Spec

what you are building, and what you are not

Decision record

MADR · context, options, consequences

Handover

what the next person needs

Changelog

Keep a Changelog 1.1.0

EXISTING EDITS · 3 PROJECTS · 128 DOCUMENTS

unreviewed first

auth.md

acme-billing / specs

GitHub App installation flow rather than an OAuth app, because the App grants permission per repository...

64% unreviewed

agent · 2h ago

0004-sync.md

acme-billing / adr

We will not use a CRDT. Merging by git three-way keeps CRLF and the missing final newline...

88% unreviewed

agent · yesterday

pricing.md

acme-billing / notes

Team \$8/user/month. Sync \$4 later. The editor is free because the category settled that.

31% unreviewed

you · 3d ago

billing.md

acme-billing / specs

One payment attempt per Indian card. ₹15,000 per transaction is an architectural constant.

12% unreviewed

agent · 4d ago

README.md

acme-billing

Billing for small teams. Runs on your repo; we never hold your documents.

all reviewed

you · 1w ago

0003-scope.md

acme-billing / adr

Out of scope for v1: invoices in currencies other than INR and USD, and any refund automation.

45% unreviewed

agent · 1w ago

onboarding.md

percent / notes

Minute zero: one button, choose a folder. No account, no email field, no tour.

all reviewed

you · 2w ago

api.md

percent / specs

REST with idempotency keys on every write. Errors are RFC 9457 problem details.

22% unreviewed

agent · 2w ago


3 projects · 128 documents · 7 unreviewed

⌘K commands · ⌘P files

The bar on every card is the product showing itself before anything is open. Templates are markdown skeletons: spec, decision record, handover, changelog.

EXHIBIT 7

S2, the editor. Tree, coloured tabs, four modes, the review tint, the AI panel collapsed



Text changed since you last reviewed it gets a light tint and nothing else. Click a span and you get when it changed, which bytes, the author from the commit trailer, and, to be honest, "prompt not recorded" for a file the agent wrote outside frontmatter. The AI panel is one collapsed line. Nothing runs until you ask, ⌘J opens it, and Settings can hide it completely.

EXHIBIT 8

S3, answer in place. A spec-kit feature, kept because it is cheap

fm

spec.mdplan.mdtasks.md

ShareSearchsm

TREE

FP+

LiveEditSplitRead

written by spec-kit via Claude Code · 11 min ago · 3 questions for you

acme-billing

specs

001-invoicing

spec.mdplan.mdtasks.md

.specify

AGENTS.md

feature: 001-invoicing

status: draft

generated-by: speckit.specify · 2026-09-06

Feature Specification: Invoicing

Functional requirements

- FR-001: System MUST let an admin create an invoice from a customer record with line items, tax and a due date.
- FR-002: System MUST render the invoice as a PDF that is byte-identical across renders for the same input.
- FR-006: System MUST authenticate users via
 - A SSO through the GitHub App – no second identity path
- FR-007: System MUST retain invoices for
 - NEEDS CLARIFICATION 7 years – Indian Companies . Answer
- FR-009: System MUST charge Indian cards
 - NEEDS CLARIFICATION once, or retry on failure?

Non-goals

Multi-currency invoicing, refunds automation, and any accounting export are out of scope for this feature.

Answering writes → A: ... into the file at that byte range, exactly where spec-kit reads it back. 1 line changed, 0 lines moved. The requirements the agent wrote stay tinted until you have read them.

Questions in this spec2

FR-007 · retention periodopen

FR-009 · card retry policyopen

FR-006 · auth methodanswered

Document outline

Functional requirements5 · 5unreviewed

Non-goals1 · 1

Document history

nowyou · answered FR-006

11mspec-kit · /speckit.specify · 2,140 bytes

Mark all readShortcuts

AI edit · #J

7 unreviewed spans · 92%2 questions open

git · +1 -1 · spec.mdLF · UTF-8

EXHIBIT 9

S4, the review panel. Catching up on a file you did not watch being written

fm

0004-sync.mdauth.mdpricing.md

ShareSearchsm

TREE

FP+

LiveEditSplitRead

acme-billing / adr / 0004-sync.md

acme-billing

specs

auth.md

adr

0003-scope.md

0004-sync.md

notes

pricing.md

title: ADR-0004 Sync without a CRDT

status: proposed

date: 2026-09-02

ADR-0004 · Sync without a CRDT

Context

Two devices edit the same markdown file. Most sync products reach for a CRDT because convergence is guaranteed. Convergence is not the property we want: CRDTs interleave concurrent edits character by character, so two people finishing the same sentence produce a merged sentence neither wrote.

Decision

Git three-way merge, with divergence surfaced as a reviewable conflict. On the same inputs it preserves Windows line endings and a missing final newline, and it stops exactly where a careful writer would want to stop.

Consequences

- Conflicts happen, and users hate conflicts. We publish the conflict rate as a shipped metric with a budget.
- We never guess which version was meant. That is the whole point.

3 UNREVIEWED IN THIS FILEJ K · A

Context · 312 bytes · sonnet-5 · 2h

Two devices edit the same markdown file. Most sync products reach for a CRDT because...

a · r · s · revert · skip

Decision · 218 bytes · sonnet-5 · 2h

Git three-way merge, with divergence surfaced as a reviewable conflict...

Consequences · 96 bytes · sonnet-5 · 2h

Conflicts happen, and users hate conflicts. We publish the conflict rate...

Provenance is a state, not a moment. These stay marked until a person accepts or reverts them — today, next week, or when a teammate opens the file.

3 unreviewed spans · 88%231 words

git · clean · mainCRLF · UTF-8

Every changed span in order. j and k to move, a, r, s to accept, revert or skip, and a counter that goes to zero. This is the state nobody else keeps.

EXHIBIT 10

S2, split. Source and rendered, the same bytes

fm

pricing.md xauth.md x

ShareSearchsm

< TREEFP+

acme-billing

- notes
- pricing.md
- specs
- adr

LiveEditSplitRead

acme-billing / notes / pricing.md · source = rendered, same bytes

title: Pricing
status: decided
tags: [pricing, obsidian, teams]

Pricing

The editor is free. Obsidian settled that: the whole editor at \$0, sync at \$4, a \$50/user/year commercial licence. Everyone priced around it.

What we charge for

- **Team** - \$8/user/month. Provenance across a shared repo, review state across people, admin.

- **Sync** - \$4/user/month, later, when it is provably safe.

The arithmetic

At \$8/seat, **114 paying seats** covers the monthly nut of ₹1.09L - about twenty small teams. Churn on a team licence is far lower than on a personal one.

title: Pricing
status: decided
tags: [pricing, obsidian, teams]

Pricing

The editor is free. Obsidian settled that: the whole editor at \$0, sync at \$4, a \$50/user/year commercial licence. Everyone priced around it.

What we charge for

- **Team** — \$8/user/month. Provenance across a shared repo, review state across people, admin.
- **Sync** — \$4/user/month, later, when it is provably safe.

The arithmetic

At \$8/seat, **114 paying seats** covers the monthly nut of ₹1.09L — about twenty small teams. Churn on a team licence is far lower than on a personal one.

2 unreviewed spans · 31%96 words

git · clean · mainLF · UTF-8

EXHIBIT 11

S11, settings. The hide-all-AI switch, and the key field that does not exist yet

fm

Settings

sm

Settings

APPEARANCE

Theme

System

Editor typeface

Google Sans Code

KEYMAP

Default

AI PROVIDER

☐

Hide all AI features. Removes every AI control from the interface. The review tint stays — it is file state, not AI

Provider

Anthropic

Your API key

.....

Stored in your OS keychain. Requests go from this machine to the provider. They do not pass through us.

Model for language work

claude-sonnet-5

Structural work — links, slugs, frontmatter fills — never calls a model.

☐

Use frontmatter's hosted AI instead (metered, hard-capped, never the default)

PROVENANCE

☒

Mark machine-written text until a person reviews it

☒

Show the unreviewed bar in the file tree

Record

byte range

prompt

model

time

session id

Written to .frontmatter/provenance.jsonl beside your documents, in your repo. Plain text. Readable without us.

EXHIBIT 12

S8, refactor preview (MVP-1). The engine, made visible

fm

Rename tag #auth → #authentication

Cancel

Apply 6 files

8 FILES · 6 CHANGE · 2 REFUSED

specs/auth.md2 hunks

specs/billing.md1

specs/api.md1

adr/0003-scope.md1

notes/pricing.md1

README.md1

notes/2026-08-12.mdrefused

archive/old-auth.mdrefused

specs/auth.md · frontmatter · bytes 41–451 of 2

4 tags: [auth, github-app, security]

4 tags: [authentication, github-app, security]

acceptskip

accept all in file

specs/auth.md · body · bytes 1,982–1,9872 of 2

38 See also [[billing]] and the

39 #auth notes from the kickoff.

39 #authentication notes from the kickoff.

acceptskip

REFUSED · 2 FILES

notes/2026-08-12.md — #auth appears inside a fenced code block at line 17; rewriting code is a guess we will not make. archive/old-auth.md — the tag sits in flow form tags: auth with no brackets; the target is ambiguous. Both files are untouched.

Every file that will change is listed before a byte moves. Accept all, per file, or per hunk. Ambiguity is refused, never guessed — and the refusal says why.

EXHIBIT 13

S12, export (MVP-2). A site or a page from the folder, never hosted by us

fmGenerate · acme-billing ×

OUTPUT

Static site · tree as navigation

One shareable page

Slide deck · H2 = slide

WHAT GOES IN

specs/ · adrl · README.md

notes/ excluded

☐ Skip documents with unreviewed machine text

THEME

Plain · one column

Docs · sidebar

Anything custom is site, css in your repo.

DEPLOY

Vercel · your account

GitHub Pages

Netlify

We never host it. Files land in . /site in your repo. Deploy runs on your account, under your terms.

acme-billing.vercel.app/specs/auth

Specs

Authentication

Billing

API

Decisions

0003 Scope

0004 Sync

Authentication

We use the GitHub App installation flow rather than an OAuth app, because the App grants permission per repository instead of across the whole account.

The installation token is scoped to the repositories the user selected and expires after one hour, which means a leaked token has a bounded blast radius.

What we never do. Ask for the repo scope. Keep a second identity path.

generated by frontmatter · every paragraph reviewed by a person · source: specs/auth.md

A site is another projection of the same files. No site editor, no theme builder — edit the markdown, regenerate.

PART V

Money, market, plan

8 Price, with the churn maths attached

EXHIBIT 14

What we charge, corrected against what teams are actually gated on

Tier	Price	What it is
Individual	Free, forever, no limits	The editor, review state, revert, answer in place, your own key. This is the distribution
Pro, team	\$4–5 per user a month. Priced against GitHub Team \$4, Obsidian Sync \$4, HackMD \$5, Confluence \$6.70	Shared-vault sync, version-history retention, shared review state, comments, MCP. Not "who reviewed what": GitHub sells that at \$4
Max	Later, per team	Document gates in CI, retention beyond a year, admin
Hosted AI	Metered, at cost plus a margin	Optional, never the default

The maths, plainly. Our monthly nut is about ₹1.09L, covered by 54–78 consulting hours. Infrastructure is about \$0 at 100 users and about \$390 a month at 10,000, because we store no documents. Support is the real cost, roughly 90 founder-hours a month at scale. I withdraw "114 seats at \$8". The feature it priced is GitHub's at \$4, and accounts under \$50 a month sit in the worst-retaining band there is: top-quartile annual gross retention 60–70% across 2,100 businesses, and 23% for AI-native products. Holding 114 seats would mean replacing 34–46 of them every year, at best. Nobody comparable publishes retention. Obsidian says on record, "we don't know what suddenly causes someone to churn." So the team tier is a hypothesis until the 20 conversations in §11 tell us seat counts and what teams are gated on.

TEST

Give away what the category has proven is free. Charge for what small teams are gated on elsewhere: seats, history, sync. Never for a feature GitHub bundles. Downgrading loses features, never files.

9 What grabs users, and how we reach them

The message is now three versions, tested in weeks 1–2 instead of one line picked by us: *See what the AI wrote. See what the agent changed without asking. Know what nobody has reviewed yet.* Reddit's top twenty "AI wrote" posts contain zero asks for authorship; the upvoted pain is the second and third. Under any of them, the proof line: *git diff after an edit shows your change and nothing else. Tested on 8,513 real files, every release.*

Lead with what we refuse, because this audience is tired of AI tooling: no plugins, no code execution, no holding your files, no lock-in, and a switch that hides every AI feature, which the incumbents only shipped after a backlash. Never say revolutionary, seamless, AI-native, or byte-preserving in a headline.

EXHIBIT 15

Channels, resized to real usage, in the order we use them

When	Channel	Size, measured	What we do
Week 2	Free Obsidian plugin, the live-preview fix	501 likes, 117 posts, open since 2022. Against it: Obsidian patches live preview monthly, and the nearest plugin-level render fix has 1,454 downloads	Ship it. Kill signal: under 200 installs in 14 days
Week 2 on	superpowers first, spec-kit second, plan-mode files third	22–30k, then about 1,660 active public projects a month. Plan-mode size untested	A README line, an MCP entry, "open the spec here" in their docs. A PR to a maintainer is an outbound act, so we ask you first
Now on	Our own writing		The research behind this document. Nobody has published it
MVP-0 exit	Show HN	Zed's scored 43, Cursor's 9	After 6 of 10 strangers keep it. One shot
MVP-1	Chrome extension, Product Hunt	1.5M browser readers	Backlinks and a spike
MVP-2	Exported pages with an honest mark	Not a channel until a deploy count exists	Withdrawn from "always"

How we use the users we get. The ten pilot strangers are the interviews we have never done. Every accept and revert is a preference signal, logged only in a local sidecar they can read, because the same architecture that keeps their documents out of our hands keeps their usage out too. The plugin's installers are the list we do not have.

10 Challenges, and what we do about each

EXHIBIT 16

Ranked by what would actually stop us. Two rows are new

Challenge	What we do
Review state turns out to be nice-to-have	Two-week test, three headline versions, kill gate: fewer than 4 of 10 call it useful unprompted
Anthropic ships the review surface. Plan-as-markdown with comments exists since 2.1.70; the diff-review request is open at 262 reactions	Persistence is what they do not build: across sessions, people and tools. If they build that too, we are a feature, and we should know by week 12
The MVP-0 exit test cannot be observed. Under our own architecture we cannot see who "keeps it". Obsidian cannot either	A written measurement decision before week 10. Opt-in, local sidecar counts they can read. Never a vendor usage database
Nobody pays for the team tier	Run the 20 conversations in weeks 1–2, not at MVP-1. Ask what they are gated on and how many seats they are
Zed or Cursor add vault semantics, one sprint	Depth in markdown, and the state they do not keep
Distribution	Three cheap channels, resized to real numbers. The private-repo multiplier is unmeasured; ask the ten strangers
The 83% refusal is a class, not one bug	Falsified if a patched corpus still refuses more than 10 of 7,969, by day 24
Every AI call bills us	Your own key before any stranger
Client work crowds out product	Capped at 78 hours a month, written down
GST	Reverse charge has no turnover floor. Registration starts with the first API purchase

11 Ninety days, and the decisions

EXHIBIT 17

Something observable every fortnight

Weeks	What happens	Outcome
0	Rotate the two access tokens. Decide the buyer. Write the web byte-path and the measurement decision	Done, written down
1–2	Four tests, build nothing . 10 developers shown the review mock in both framings (who wrote it, versus changed since you reviewed) and asked which toolchain they use. Plugin shipped. 5 people who bill for documents asked about approval and sign-off, not provenance. 3 landing pages, one per headline. 20 team conversations on seats and gating	Kill signals read at day 14
3	Go or no-go	A written decision either way
3–4	CI with a deliberate red run. NF-3, then NF-1. Prototype content-hash review state on one spec-kit repo and one superpowers repo	CI fails a broken commit. Refusals fall from 83%. The tint works on files we did not write
5–7	Review state store and rendering. Your own key. Structured output	A stranger's key, a stranger's repo, tinted spans
8–9	Revert, the review panel, answer in place, engine wired behind every write	The fifteen-second demo from §1
10–12	Table stakes. Ten strangers with their own repos and agents	6 of 10 keep it, by the measure written in week 0

EXHIBIT 18

The decisions

#	Decision	By
1	Accept the change of direction. Review state is the product, authorship is a detail, answer in place is a feature, the team tier is a hypothesis	The meeting
2	Which buyer first: the developer (reachable, but AI-native products under \$50 keep 23% of revenue a year) or the person who is liable when a document is wrong	Week 1
3	Share. A hosted read-only viewer (our old record's assumption, and a hosting obligation) or permalink plus export (stateless, but nothing for a plain folder or a private repo)	Week 2
4	The two unrotated access tokens. An action, not a decision	Today

RISK

What would make me say stop. Fewer than 4 of 10 developers finding review state useful under either framing, and the 20 team conversations turning up nothing they are gated on that GitHub does not already sell. Then there is no product here, only an engine, a services business and a free plugin people like. Not a failure. A smaller, truer version of the same work, and it pays sooner.

PART VI

The rest, in one page

12 The engine: what exists, what is broken, what needs work

EXHIBIT 19

Built, broken, to be optimised

Built	25,407 lines, 1,575 tests. 8,513 real files from strangers' vaults, byte-pinned, zero corruption. <code>git diff</code> after an edit shows only the edit, tested every release. Six AI routes, a five-provider chain, idea mode
Broken	NF-1: a column-zero list item in frontmatter refuses 83% of real vaults. 4 days. NF-3: a bare-CR fence adds a second frontmatter block. 3 days. No CI , and four gates have reported green while blind. 1 day. Engine unwired: one symbol from one of thirteen files reaches product code. 8 days. Byte budget is an echo. 2 days
Optimise	Incremental parsing , a vendored @Lezer/markdown, so review spans survive every keystroke cheaply. The offset map , bytes to UTF-16, our riskiest seam; crossing it unmapped corrupts Chinese, Japanese, Korean and Arabic text silently. Search: today the whole vault ships to the client, 77 MB parsed per cold start; move to server full-text, client only for the open file. Refusal and conflict rates published, with budgets

What we publish	Time to first keystroke, cold. Typing latency on a 10,000-word document. Corpus refusal rate (83% to near zero). Merge conflict rate
-----------------	--

13

What we render, and the markdown rules

EXHIBIT 20

Custom rendering, and what we never do to a file

Modes	Live, Edit, Split, Read. Nested constructs inside list items render properly: the 501-like bug, shipped as the free plugin first
We render	Callouts > [!kind]. Mermaid, rendered, not extended. MADR and Nygard decision records. RFC 7322. Keep a Changelog 1.1.0. Frontmatter as a properties panel
We do not pretend	Runbooks and PRDs have no standard body. We offer a structure and say so. Every "PRD standard" claim out there is unsourced
The rule	We add nothing to markdown that breaks it somewhere else. A file we touch still renders on GitHub, in Obsidian, in a plain editor, next year
How	Prose annotations as callouts, which have no closing marker to lose; an unclosed fence swallows the rest of the document. Tags and links written back in the shape the file already uses, never converted. Wikilinks and normal links both read. Review state and comments in a plain-text sidecar next to the file, readable without us
What we join	Read AGENTS.md and CLAUDE.md, never compete with them. Be an MCP server the user's agent connects to. We are a git client, never our own versioning. Obsidian conventions preserved

14

Where we stand

EXHIBIT 21

Strengths, weaknesses, opportunities, threats

Strengths	The engine is real and correct where the incumbents admit they are not. 88.6% of a shipping editor exists. Six AI routes live. Two founders funded by services, not a clock
Weaknesses	Zero users, zero interviews. No CI, no BYO key, no structured output. 83% of vaults refused today. Every channel is borrowed. No evidence yet for the paid tier. Three published numbers failed a second check
Opportunities	Tens of thousands of active projects generating specs with nowhere to review them. Review state is the one thing no incumbent persists. A 501-like bug we can fix in a week. A hide-all-AI switch the incumbents shipped only after backlash
Threats	Anthropic ships persistent review. Plan-as-markdown exists since 2.1.70 and diff review is open at 262 reactions. Zed or Cursor add vault semantics in a sprint. Obsidian ships first-party AI editing. We run out of attention before revenue

15

Operating rules

EXHIBIT 22

The lines we hold

D2C and B2B	Free to build the audience, paid to build revenue. No enterprise feature until a customer refuses to pay without it. SSO at about 20 seats when asked, SOC 2 above about 50 or for a regulated buyer, a merchant of record for EU and US procurement
How users hear from us	We collect an email only when there is a reason to: a payment, an invite, a recovery request. Never at first run. Breaking changes get two weeks' notice in the product. Price changes are emailed before the next charge. Every notice also lives in a quiet in-product list, so "we told you" is something they can check
How they reach us	One address, a published response window, and the free tier has no service commitment, said plainly on the page
What we refuse to measure	Session replay, keystroke telemetry, document content. The DOM would be their private document. Usage counts live in a local sidecar they can read
Mobile	About 15% of complaints are mobile. Not in v1. When it comes, a reader and reviewer, never a byte-exact editor on a phone keyboard
Accessibility, non-English	WCAG 2.2 AA, all or nothing. Our body-faint token fails at 1.984:1 and gets fixed in MVP-0. The tint is never colour alone; the panel and the keyboard carry the same information. CJK: the byte map, a bigram search tokeniser, and a written decision on v1 scope
Security	No plugins, no code execution: the biggest single protection we have. The AI never gets ambient repo access. A malformed repo is refused, not repaired. Reverse charge under CGST §24(iii) has no turnover floor

16 What I think

- The engine is the best thing either of us has built. Two rounds of checking did not touch it.
- We were wrong about the pitch three times: bytes, then authorship, then a generator. Each time the evidence was public and we had not opened it. What survives is narrower. A state nobody keeps, on files everyone now generates, in a place nobody owns. It is a product only if ten strangers say so in two weeks.
- The open question is not distribution any more. It is whether anyone pays. We have a free product with a channel and no evidence for the paid tier. The twenty conversations settle it, in week one.