


frontmatter

Product and technical plan: pilot, pro, max

Technical documentation, Studio Zephyrus. Market figures re-checked against 157 primary sources on 2026-09-06; browser support, repo-tooling and hosting figures added 2026-09-08. Corrections to earlier versions are marked where they apply. Features are labelled by their evidence: **public evidence**, **founder user test**, or **untested**.

KEY FIGURES



PART I


Product

1 Definition

GOOD

frontmatter is a markdown editor that shows what an agent changed and what nobody has reviewed yet, on files that stay in the user's own repo. Every span changed since a person last read it is tinted. One key reverts any span. Author and prompt are shown when the write carried them. The file is edited byte-exactly: only the requested range changes.

Core loop



Not in the product: no plugins or code execution, no hosting of user content, no vendor database holding document bytes, no project management, no chat sidebar, no mobile editor in v1.

Demo, fifteen seconds. Run the agent's spec command. Open the spec in frontmatter: every paragraph tinted. Read one, it clears. Revert one, `git diff` shows only that line. The panel lists what is still unread.

2 Feature specification

Each feature: what it does, how it works, what it needs, when it ships, and whether it works online without an install and offline in the desktop app. Days are engineering estimates for the current team; they assume the engine fixes in §9 land first.

EXHIBIT 1

Features, specified

#	Feature	What it does	How it works	Needs	Ships	Web, no install	Desktop, offline	Evidence
F1	Editor	Live, Edit, Split, Read. Tree, tabs, quick-switch, palette, search, properties panel for frontmatter	CodeMirror 6 over the splice engine; every write is a byte-range replacement	Exists. NF-1/NF-3 fixes (7d), engine wiring (8d)	MVP-0	Yes	Yes	Public: editor is table stakes

#	Feature	What it does	How it works	Needs	Ships	Web, no install	Desktop, offline	Evidence
F2	Review state	Tints every span changed since a person last read it; one-key revert; review panel with a counter	Per-span content hash in <code>.frontmatter/review.jsonl</code> beside the files, in the repo. Spans anchored by byte range + hash, relocated by the splice locator when the file changes	Store 6d, rendering 5d, revert 3d, panel 4d	MVP-0	Yes	Yes	Public: unrequested-change pain (262 reactions on diff review). Untested as a product — the two-week test
F3	Attribution	Author, model, time and prompt on hover	Read from what the write carried: an MCP write (F11), a commit trailer, a git-ai note, signed provenance metadata. Absent when none exists, and the card says so	Reader for trailers/notes 3d	MVP-0 (readers), MVP-1 (MCP)	Yes	Yes	Public: authorship marking alone did not sell (iA Writer 7)
F4	Answer in place	Answers [NEEDS CLARIFICATION] and Q: → A: markers inline; one line changes	Marker detection + splice into the marker's byte range	3d	MVP-0	Yes	Yes	Public: ~1,500 live public files, spec-kit only
F5	Idea mode	Generates a PRD, FRD, BRD, product note or spec from a typed idea or from the repo (README + specs)	Exists: generate-document, five prompt kinds, <code>/api/ai/generate-doc</code> . Output arrives as tinted spans; nothing written until kept	Structured output in the port 2d; "from repo" input 2d	MVP-0	Yes, with the user's key	Needs network for the model; offline: no	Founder user test (reported). Public: standalone PRD generators have no traction — top repo 57 stars of 512
F6	Decision flow	Before generating, asks the questions the idea leaves open; answers are written into the frontmatter and shape the document	Structured question list from the LLM port; answers stored as frontmatter keys the generator reads	4d, after F5's structured output	MVP-0, inside F5	Yes, with key	No	Founder user test (reported). Public: <code>/speckit.clarify</code> , Kiro and Claude Code plan mode ship the same step free — the difference is that ours writes into the document and stays reviewable
F7	Repo docs scan	Reads every markdown file in a repo and proposes: stale sections (doc older than the code it cites), broken links, missing standard docs (AGENTS.md, CHANGELOG), and offers to generate them	No model for the scan: git dates per file, <code>package.json</code> scripts, link resolution (<code>link-doctor</code> exists). Generation reuses F5. Shows exactly what was read	Staleness heuristic 4d, missing-doc templates 2d, S13 UI 4d	MVP-0	Yes: scan runs in the browser on fetched bytes	Yes: scan needs no network	Founder direction. Public: Vale 6,087 stars and markdownlint 6,326 are the incumbents for lint; no incumbent proposes edits with review state. README generators: 11,131 stars, abandoned 2022
F8	File-scoped AI	Two verbs on the open file: fix or critique this selection; summarise this file. Collapsed by default; a hide-all switch	Bring-your-own key, request from the user's machine to the provider. Proposals arrive as tinted spans	BYO key UI + keychain 3d	MVP-0	Yes, with key	No	Public: editing supplied text is 10.6% of ChatGPT messages; incumbents hide AI until invoked
F9	Vault-wide refactor	Rename a tag, heading or key across the vault; every hunk reviewable; ambiguity refused	Engine splice across files; refusal on fenced or ambiguous matches	8d	MVP-1	Yes	Yes	Public: 86 likes on broken-links-on-rename
F10	Shared review state, comments	What each person has read, across a team; comments anchored to spans	Same sidecar, per-person entries; comments as a sidecar too, never in the file (every in-file syntax renders as garbage on GitHub)	10d	MVP-1	Yes	Local only until push	Public: GitHub Team sells reviewers at \$4; untested whether teams pay for this

#	Feature	What it does	How it works	Needs	Ships	Web, no install	Desktop, offline	Evidence
F11	MCP server	The user's agent writes through frontmatter, so writes carry the prompt	Local MCP server in the desktop app; agents connect by config	6d	MVP-1	No: needs a local process	Yes	Untested whether agents route writes through a third-party tool
F12	Sync	Multi-device, provably safe	Git merge, never a CRDT; conflicts surfaced, never guessed	15d+	MVP-2	Yes	Local edits, sync on reconnect	Public: sync is the one proven individual price
F13	Export	Single document or folder to HTML, PDF, copy-as-HTML; decks only if a pilot user asks	Templating, no model	5d	MVP-2	Yes	Yes	Public: export beats decks 2:1; sites are wanted hosted, which is refused

Founder-validated features. F5 and F6 were reported as tested with users by the founders. Public evidence for them is weak as standalone tools and does not contradict an in-editor test. They ship in MVP-0 on that basis, inside the review-state loop: generated text arrives tinted and is not written until kept.

3 Access model

Two ways to use frontmatter. What each can do is fixed by the browser and the platform, not by the plan.




EXHIBIT 2
What works where


Capability	Web, GitHub repo	Web, local folder (Chrome, Edge)	Web, Firefox, Safari, mobile	Desktop app
Open and edit files	Yes; writes become commits or PRs	Yes, direct writes with a per-folder permission	Repo connect or drag-in only	Yes
Review state (F2)	Yes; sidecar committed with the files	Yes	Repo connect only	Yes
Repo docs scan (F7)	Yes, on fetched bytes	Yes	Repo connect only	Yes, offline
Idea mode, AI panel (F5, F6, F8)	With the user's key	With key	With key	With key; needs network for the model
MCP server (F11)	No	No	No	Yes
Offline	Cached open files only	Cached	No	Everything except the model
Local git, watcher, index	No	No	No	Yes

Browser support, fetched 2026-09-08 (caniuse): the File System Access API is supported in desktop Chrome and Edge and in none of Firefox, Safari, iOS Safari, Android Chrome or Samsung Internet. **GitHub's REST API supports CORS**, so a repo's bytes can be fetched in the browser with the user's token; the vendor server never sees them.

What the user downloads. A Tauri v2 desktop app: macOS .dmg (notarised, Apple Developer ID \$99/year), Windows .exe (code-signed, \$150–400/year, hardware token required), Linux Applmage. It contains the editor, the engine, a git client, the MCP server and the search index. It contains no model. Tauri reuses the operating system's webview instead of bundling a browser; the installer size is measured at the first build. The auto-updater must be signed and configured before the first public build. Note: the current src-tauri config still carries the sibling app's identity (productName sgnk-md, ai.sgnk.md) and must be re-identified.

If the user does not download. They get F1–F8 in the browser against a GitHub repo, or against a local folder on Chrome and Edge. They do not get MCP, local git, a background watcher, or offline beyond cached files. The web app is the trial; the desktop app is the daily surface. Both are one build with a filesystem adapter, so there is one product to maintain.

4 Architecture



Data path. Document bytes move between the editor, the engine and the user's git. In the web app they are fetched from GitHub in the browser. The control plane holds identity, team membership, entitlements and billing, and no document bytes. The AI request goes from the user's machine to their provider. Hosted AI, if ever enabled, is metered and hard-capped and never the default.

Sidecar format. `.frontmatter/review.jsonl`, one JSON line per span: file path, byte start, byte end, content hash, reviewed-by, reviewed-at, and, when known, author, model, prompt. Plain text, documented, readable without frontmatter. Spans are relocated by the splice locator on every edit; the journal is not the anchor, because git, other editors and agents bypass it. The nearest shipped pattern is SpecKit Companion's `.spec-context.json` (9,367 installs); frontmatter's must be visibly better: per-span, per-person, any tool.

Cost rules. Nothing runs in the background; spend follows use. Output tokens are 75.8% of model spend, so output is capped, not input. Infrastructure is about \$0 at 100 users and about \$390/month at 10,000 because no documents are stored. Support is the real cost at scale, about 90 founder-hours a month at 10,000 users.

PART II

Market

5 Who does what, and where each stops

EXHIBIT 3

Fetches 2026-09-06 and 2026-09-08

	Tool	Reach	What it does	Where it stops
Spec writers	obra/superpowers	75–85k public repos; 22–30k active in 30 days; marketplace counter 1,009,371 (cached)	Brainstorm, design, write the spec, commit it	Terminal. No editor, no review state
	github/spec-kit	11,072 public repos; about 1,660 active	<code>/speckit.specify</code> , <code>/speckit.clarify</code> (in chat), plan, tasks, implement	Writes markdown, then leaves

	Tool	Reach	What it does	Where it stops
	Plan mode, Claude Code and Cursor	Out-scores both named tools on Reddit (1,776 / 1,409 / 1,329 points vs 38 or fewer)	Plans as markdown files	Where the files live is untested
Review surfaces	Claude Code, VS Code extension	24.89M installs	Plans as a markdown document with inline comments since 2.1.70 (2026-03-06)	Per session. Diff-review UI open at 262 reactions
	git-ai; VS Code 1.118	61,517 installs; AI co-author trailers by default	Line attribution after the fact; commit trailers	Commit granularity, nothing per span
	SpecKit Companion	9,367 installs	Review state in <code>.spec-context.json</code> per spec	Spec-kit only
Docs tooling	Vale, markdownlint	6,087 and 6,326 stars	Prose and markdown lint	Lint only; no proposals, no review state
	readme-md-generator, readme-ai	11,131 (abandoned 2022), 2,980	Generate a README	One file, no repo context
	PRD generators	512 repos on GitHub; the top has 57 stars	Generate a PRD from a prompt	No traction as standalone tools
Editors	Obsidian	Free; Sync \$4; Publish \$8/site; Commercial \$50/yr	AI via plugins: Copilot 1.83M downloads, Claudian 2.02M	None records what changed or was reviewed
	Cursor, Zed	~\$20; free	Agent editing for code; Zed reviews per hunk	Whole-file rewrite; no vault semantics
Browser	Obsidian Web Clipper, Markdown Viewer	1,000,000 and 500,000 Chrome users	Capture and read markdown	A distribution channel
Hosted sites	mdown.ai, Obsidian Publish	Non-engineers; \$8/site/month	Markdown to a hosted website	The host role frontmatter refuses

Prices. The editor is worth ₹0 in this category. The one proven individual price is sync, about \$4. Obsidian's \$50/year commercial licence confers no functional benefit, became optional in February 2025, and is bought by 25+ seat organisations; it is not evidence of what a five-person team pays.

6 Vocabulary audit

EXHIBIT 4

Terms from earlier plans, checked

Term	Finding	Status
Context packs, handover generation	Free from GitHub and Anthropic; the "median 2 points" graveyard was the Hacker News baseline	Cut
Decision-flow renders	0.30 of the median plugin, wrong buyer	Cut
Kickoff prompt, vendor CDN	Worse than a slash command; breaks "no document bytes on our server"	Cut
Self-improving AIOS	One feedback label from 6,884 decisions	Internal tooling only
Local LLM fallback	Five cloud free tiers, not local; constrained small models fabricate in 10 of 13	Deferred
Byte-exactness as the headline	4 complaints in 12,556	Mechanism, not pitch
"See what the AI wrote"	Shipped by iA Writer 7 in 2023; 208 downloads for the Obsidian port; best VS Code extension 366 installs; one Reddit ask	Attribution is shown when available, not the headline
"496,000 stars" as channel size	Stars measure attention to a CLI	Replaced by active-repo counts
"Uncopyable"	Trailers and git-ai attribute at commit level	Struck
"62,976 unanswered markers"	49,408 were spec-kit's own checklist line	About 1,500 live markers

7 Evidence status per feature

EXHIBIT 5

What is proven, what is reported, what is untested

Status	Features	What settles it
Public evidence	Editor table stakes (F1), unrequested-change pain (F2 motive), attribution-alone does not sell (F3), marker counts (F4), AI hidden until invoked (F8), sync as a price (F12), export over decks (F13)	Done
Founder user test, reported	Idea mode (F5), decision flow (F6), the docs-scan direction (F7)	The founders' test notes, attached to the plan; then the ten-stranger pilot
Untested	Review state as a product (F2), teams paying for shared review (F10), agents writing through MCP (F11), the web byte-path decision, the MVP-0 exit measure	Weeks 1–2 tests (\$14)

PART III

Delivery

8 Tiers

EXHIBIT 6

MVP-0 pilot, MVP-1 pro, MVP-2 max

	MVP-0 pilot, free forever	MVP-1 pro, teams	MVP-2 max
Engine	NF-1, NF-3 fixed; CI with a deliberate red run; engine wired behind every write; real byte budget	Vault-wide refactor (F9)	Sync (F12)
Editor	F1, all modes and navigation	Conflict view	Multi-device
Review	F2 review state; F3 attribution from trailers and notes; F4 answer in place	F10 shared state, comments; F3 with prompts via F11 MCP	Version-history retention
Generation	F5 idea mode; F6 decision flow; F7 repo docs scan and missing-doc generation		F13 export
AI	F8 panel, BYO key, two verbs, hide-all switch; structured output in the port	Hosted AI, metered, capped, off by default	Document gates in CI
Surfaces	Web app (repo, local folder on Chrome/Edge) and desktop app, one build	Chrome extension; MCP server	
Channels	Free Obsidian live-preview plugin, week 2; kill: under 200 installs in 14 days	superpowers and spec-kit docs; plan-mode files	
Absent	Sign-up, billing, sync, mobile, hosting, hosted AI, local model	Mobile, plugins, chat sidebar	Plugins
Exit test	6 of 10 strangers still using it after two weeks, by the measure written in week 0	A team pays within 20 conversations	Corpus passes byte-identical after sync

9 Technical plan

Fix first. These block everything else.

#	Problem	Effect	Days
1	NF-1: column-zero list item in frontmatter	83% of real vaults refused	4
2	NF-3: bare-CR frontmatter fence	A second frontmatter block is added	3
3	No CI; four gates have reported green while blind	Shipped on trust	1
4	Engine unwired: one symbol from one of thirteen files reaches product code	Product does not use the engine	8
5	Byte budget is an echo	No correctness gate	2
6	No bring-your-own key; every AI request bills the operator	Cannot expose AI to a stranger	3
7	LLM port returns a bare string; JSON is sliced between brackets	No structured output for F5, F6, F7	2

#	Problem	Effect	Days
8	Web byte path undecided	"Zero document bytes on our server" is unverified for the web app	Decision
9	src-tauri carries the sibling app's identity	Desktop build ships as sgnk-md	1

Sequence. Fix-first (24 days) → review state F2 (18 days) → attribution readers F3 and answer in place F4 (6 days) → idea mode with structured output and decision flow F5, F6 (8 days) → repo docs scan F7 (10 days) → AI panel F8 (3 days) → table stakes and polish. About 70 engineering days for MVP-0, or ten weeks at the team's real availability; the earlier 41-day figure assumed full-time work and was 2.5× optimistic.

Resources. Two founders and the existing team. Infrastructure at pilot scale: free tiers (Vercel, Postgres, R2). Model cost: zero, all AI on the user's key. Certificates: \$99/year Apple, \$150–400/year Windows plus a hardware token. No new hires for MVP-0. The pilot's ten users bring their own repos and their own agents.

Engine work beyond the fixes. Vendored @lezer/markdown for incremental parsing, so review spans survive every keystroke cheaply. The byte-to-UTF-16 offset map, the highest-risk seam: an unmapped crossing corrupts Chinese, Japanese, Korean and Arabic text silently. Search: the whole vault currently ships to the client, 77 MB parsed per cold start; move to server full-text search in the web app and a local index on desktop. Published measures: time to first keystroke cold, typing latency on a 10,000-word document, corpus refusal rate, merge conflict rate.

10 Screens

Built from the two hand-drawn layouts. Clickable prototype, public: [frontmatter-prototype-sagnik.vercel.app](#). Same file in the repo at docs/prototype/frontmatter-prototype.html. Screens are on the strip at the bottom of the page.

EXHIBIT 7

S0 launcher: blank, four templates, existing edits sorted unreviewed-first

fm

Search every document, every project

Open folder

New

START SOMETHING

Blank document

empty file · yours from byte one

Spec

what you are building, and what you are not

Decision record

MADR · context, options, consequences

Handover

what the next person needs

Changelog

Keep a Changelog 1.1.0

EXISTING EDITS · 3 PROJECTS · 128 DOCUMENTS

unreviewed first

auth.md

acme-billing / specs

GitHub App installation flow rather than an OAuth app, because the App grants permission per repository...

64% unreviewed

agent · 2h ago

0004-sync.md

acme-billing / adr

We will not use a CRDT. Merging by git three-way keeps CRLF and the missing final newline...

88% unreviewed

agent · yesterday

pricing.md

acme-billing / notes

Team \$8/user/month. Sync \$4 later. The editor is free because the category settled that.

31% unreviewed

you · 3d ago

billing.md

acme-billing / specs

One payment attempt per Indian card. ₹15,000 per transaction is an architectural constant.

12% unreviewed

agent · 4d ago

README.md

acme-billing

Billing for small teams. Runs on your repo; we never hold your documents.

all reviewed

you · 1w ago

0003-scope.md

acme-billing / adr

Out of scope for v1: invoices in currencies other than INR and USD, and any refund automation.

45% unreviewed

agent · 1w ago

onboarding.md

percent / notes

Minute zero: one button, choose a folder. No account, no email field, no tour.

all reviewed

you · 2w ago

api.md

percent / specs

REST with idempotency keys on every write. Errors are RFC 9457 problem details.

22% unreviewed


agent · 2w ago

3 projects · 128 documents · 7 unreviewed

⌘ commands · ⌘P files

EXHIBIT 8

S2 editor: tree, tabs, four modes, review tint, AI panel collapsed



Changed-since-reviewed spans carry a light tint and nothing else. The hover card shows when it changed, which bytes, the author from the commit trailer, and "prompt not recorded" when the write did not carry one. The AI panel is one collapsed line; nothing runs until asked.

EXHIBIT 9

S13 repo docs scan: stale, broken, missing, and what was read to decide

fm

Docs scan · acme-billing ×

auth.md ×

🔍 Rescan

✓ Review 6 proposals

MARKDOWN IN THIS REPO · 41 FILES

README.md

AGENTS.md

CHANGELOG.md

docs 12

api.md

install.md

architecture.md

specs 9

auth.md

billing.md

adr 4

notes 14

PROPOSALS · 6

scanned 41 .md · package.json · git log · 0 source files read

STALE · README.md \$ Install

package.json changed 2026-08-30 · README last touched 2026-06-12

14 npm run dev # starts on :3000

14 npm run dev # starts on :3100 (see package.json "dev")

open as proposal dismiss

EVIDENCE: SCRIPTS.DEV IN PACKAGE.JSON

BROKEN · docs/api.md · 3 links

link-doctor

22 See [[billing-v1]] for retries.

22 See [[billing]] for retries.

open as proposal dismiss

TARGET RENAMED 2026-08-02

MISSING · AGENTS.md

from tree + README + specs/

Generate a starter: stack, run commands, where specs live, what never to touch.

generate arrives tinted skip

YOUR KEY · CLAUDE-SONNET-5

DOCUMENT · from an idea or the repo

idea mode

PRD · FRD · BRD · product note · spec – from a sentence you type, or from README + specs/.

generate PRD FRD BRD

ARRIVES AS TINTED SPANS, NOTHING WRITTEN UNTIL YOU KEEP IT

What was read to make these: 41 markdown files, package.json, git log for dates. No source file contents.

Everything stays on this machine (desktop) or in your browser (web, fetched with your GitHub token). Nothing is sent to us.

What we read

41 markdown files, full text

1 package.json · scripts and deps only

Log git dates per file, no diffs

0 source files. Never, unless you point at one

Source

Desktop: local folder ~/code/acme-billing. Web: GitHub repo, fetched in your browser with your token. Our server sees no document bytes.

When it runs

Only when you press Rescan. Nothing in the background. Generation uses your key; scanning uses no model at all.

Exclude folders

Shortcuts

6 proposals 2 stale · 1 broken · 2 missing · 1 offer

read-only until you accept

Every proposal names its evidence. The right rail lists exactly what was read: 41 markdown files, package . j son scripts, git dates, zero source files. Scanning uses no model; generation uses the user's key.

EXHIBIT 10

S3 answer in place

fm

spec.mdplan.mdtasks.md

ShareSearchsm

TREE

acme-billing

specs

001-invoicing

spec.mdplan.mdtasks.md

.specify

AGENTS.md

LiveEditSplitRead

written by spec-kit via Claude Code · 11 min ago · 3 questions for you

feature: 001-invoicing
status: draft
generated-by: speckit.specify · 2026-09-06

Feature Specification: Invoicing

Functional requirements

- FR-001: System MUST let an admin create an invoice from a customer record with line items, tax and a due date.
- FR-002: System MUST render the invoice as a PDF that is byte-identical across renders for the same input.
- FR-006: System MUST authenticate users via

A SSO through the GitHub App – no second identity path
- FR-007: System MUST retain invoices for

NEEDS CLARIFICATION 7 years – Indian Companies . Answer
- FR-009: System MUST charge Indian cards

NEEDS CLARIFICATION once, or retry on failure?

Non-goals

Multi-currency invoicing, refunds automation, and any accounting export are out of scope for this feature.

Answering writes → A: _ into the file at that byte range, exactly where spec-kit reads it back. 1 line changed, 0 lines moved. The requirements the agent wrote stay tinted until you have read them.

7 unreviewed spans · 92%2 questions opengit · +1 -1 · spec.mdLF · UTF-8

Questions in this spec2

FR-007 · retention periodopen

FR-009 · card retry policyopen

FR-006 · auth methodanswered

Document outline

Functional requirements5 · 5unreviewed

Non-goals1 · 1

Document history

nowyou · answered FR-006

11mspeckit · /speckit.specify · 2,140 bytes

Mark all readShortcuts

AI edit · #J

EXHIBIT 11

S4 review panel

fm

0004-sync.mdauth.mdpricing.md

ShareSearchsm

TREE

acme-billing

specs

auth.md

adr

0003-scope.md

0004-sync.md

notes

pricing.md

LiveEditSplitRead

acme-billing / adr / 0004-sync.md

title: ADR-0004 Sync without a CRDT
status: proposed
date: 2026-09-02

ADR-0004 · Sync without a CRDT

Context

Two devices edit the same markdown file. Most sync products reach for a CRDT because convergence is guaranteed. Convergence is not the property we want: CRDTs interleave concurrent edits character by character, so two people finishing the same sentence produce a merged sentence neither wrote.

Decision

Git three-way merge, with divergence surfaced as a reviewable conflict. On the same inputs it preserves Windows line endings and a missing final newline, and it stops exactly where a careful writer would want to stop.

Consequences

- Conflicts happen, and users hate conflicts. We publish the conflict rate as a shipped metric with a budget.
- We never guess which version was meant. That is the whole point.

3 unreviewed spans · 88%231 wordsgit · clean · mainCRLF · UTF-8

3 UNREVIEWED IN THIS FILE

Context · 312 bytes · sonnet-5 · 2h

Two devices edit the same markdown file. Most sync products reach for a CRDT because...

a · r · s · revert · skip

Decision · 218 bytes · sonnet-5 · 2h

Git three-way merge, with divergence surfaced as a reviewable conflict...

Consequences · 96 bytes · sonnet-5 · 2h

Conflicts happen, and users hate conflicts. We publish the conflict rate...

Provenance is a state, not a moment. These stay marked until a person accepts or reverts them — today, next week, or when a teammate opens the file.

EXHIBIT 12

S2 split: source and rendered, the same bytes

pricing.md

auth.md

Share

Search

sm

acme-billing

notes

pricing.md

specs

adr

Live

Edit

Split

Read

acme-billing / notes / pricing.md · source = rendered, same bytes

title: Pricing

status: decided

tags: [pricing, obsidian, teams]

Pricing

The editor is free. Obsidian settled that: the whole editor at \$0, sync at \$4, a \$50/user/year commercial licence. Everyone priced around it.

What we charge for

- **Team** - \$8/user/month. Provenance across a shared repo, review state across people, admin.

- **Sync** - \$4/user/month, later, when it is provably safe.

The arithmetic

At \$8/seat, **114 paying seats** covers the monthly nut of ₹1.09L - about twenty small teams. Churn on a team licence is far lower than on a personal one.

title: Pricing

status: decided

tags: [pricing, obsidian, teams]

Pricing

The editor is free. Obsidian settled that: the whole editor at \$0, sync at \$4, a \$50/user/year commercial licence. Everyone priced around it.

What we charge for

• Team — \$8/user/month. Provenance across a shared repo, review state across people, admin.

• Sync — \$4/user/month, later, when it is provably safe.

The arithmetic

At \$8/seat, **114 paying seats** covers the monthly nut of ₹1.09L — about twenty small teams. Churn on a team licence is far lower than on a personal one.

2 unreviewed spans · 31% 96 words

git · clean · main LF · UTF-8

EXHIBIT 13

S11 settings: hide-all-AI switch, the key field that does not exist yet

Settings

sm

Settings

APPEARANCE

Theme

System

Editor typeface

Google Sans Code

KEYMAP

Default

AI PROVIDER

☐ Hide all AI features. Removes every AI control from the interface. The review tint stays — it is file state, not AI

Provider

Anthropic

Your API key

.....

Stored in your OS keychain. Requests go from this machine to the provider. They do not pass through us.

Model for language work

claude-sonnet-5

Structural work — links, slugs, frontmatter fills — never calls a model.

☐ Use frontmatter's hosted AI instead (metered, hard-capped, never the default)

PROVENANCE

☒ Mark machine-written text until a person reviews it

☒ Show the unreviewed bar in the file tree

Record

byte range

prompt

model

time

session id

Written to .frontmatter/provenance.jsonl beside your documents, in your repo. Plain text. Readable without us.

EXHIBIT 14

S8 refactor preview, MVP-1

fm

Rename tag #auth to #authentication

Cancel

Apply 6 files

8 FILES · 6 CHANGE · 2 REFUSED

specs/auth.md2 hunks

specs/billing.md1

specs/api.md1

adr/0003-scope.md1

notes/pricing.md1

README.md1

notes/2026-08-12.mdrefused

archive/old-auth.mdrefused

specs/auth.md · frontmatter · bytes 41-451 of 2

4 tags: [auth, github-app, security]

4 tags: [authentication, github-app, security]

acceptskipaccept all in file

specs/auth.md · body · bytes 1,902-1,9072 of 2

38 See also [[billing]] and the

39 #auth notes from the kickoff.

39 #authentication notes from the kickoff.

acceptskip

REFUSED · 2 FILES

notes/2026-08-12.md — #auth appears inside a fenced code block at line 17; rewriting code is a guess we will not make. archive/old-auth.md — the tag sits in flow form tags: auth with no brackets; the target is ambiguous. Both files are untouched.

Every file that will change is listed before a byte moves. Accept all, per file, or per hunk. Ambiguity is refused, never guessed — and the refusal says why.

EXHIBIT 15

S12 export, MVP-2

fm

Generate · acme-billing

Write to /site

Deploy to Vercel

OUTPUT

Static site · tree as navigation

One shareable page

Slide deck · H2 = slide

WHAT GOES IN

specs/ · adr/ · README.md

notes/ excluded

☐ Skip documents with unreviewed machine text

THEME

Plain · one column

Docs · sidebar

Anything custom is site.css in your repo.

DEPLOY

Vercel · your account

GitHub Pages

Netlify

We never host it. Files land in ./site in your repo. Deploy runs on your account, under your terms.

acme-billing.vercel.app/specs/auth

Specs

Authentication

Billing

API

Decisions

0003 Scope

0004 Sync

Authentication

We use the GitHub App installation flow rather than an OAuth app, because the App grants permission per repository instead of across the whole account.

The installation token is scoped to the repositories the user selected and expires after one hour, which means a leaked token has a bounded blast radius.

What we never do. Ask for the repo scope. Keep a second identity path.

generated by frontmatter · every paragraph reviewed by a person · source: specs/auth.md

A site is another projection of the same files. No site editor, no theme builder — edit the markdown, regenerate.

Commercial

11 Pricing

EXHIBIT 16

Tiers and prices

Tier	Price	Contents
Individual	Free, no limits	F1–F8, own key
Pro, team	\$4–5 per user per month, priced against GitHub Team \$4, Obsidian Sync \$4, HackMD \$5, Confluence \$6.70	Shared-vault sync, version-history retention, shared review state, comments, MCP. Not "who reviewed what" alone: GitHub sells reviewers at \$4
Max	Per team, later	Document gates in CI, retention beyond a year, admin
Hosted AI	Metered at cost plus margin	Optional, off by default

Retention. Accounts under \$50/month sit in the worst-retaining band: top-quartile annual gross retention 60–70% (ChartMogul, 2,100 businesses), 23% for AI-native products under \$50. No comparable editor publishes retention; Obsidian states on record that it does not measure churn. The earlier "114 seats at \$8 covers the nut" is withdrawn: the feature was GitHub's at \$4 and the band would require replacing 34–46 of 114 seats a year. **The paid tier is a hypothesis until 20 team conversations return seat counts and what teams are gated on.** Monthly nut: about ₹1.09L, covered by 54–78 consulting hours.

12 Go-to-market

Headline, three variants tested in weeks 1–2: *See what the AI wrote · See what the agent changed without asking · Know what nobody has reviewed yet.* Reddit's top twenty "AI wrote" posts contain no request for authorship marking; the upvoted pain matches the second and third. Proof line under any of them: `git diff` after an edit shows the change and nothing else; tested on 8,513 real files every release.

Positioning. Lead with refusals: no plugins, no code execution, no holding files, no lock-in, and a hide-all-AI switch, which Zed, VS Code and Telegram each shipped only after 412-, 30- and 181-reaction issues.

EXHIBIT 17

Channels, in order

When	Channel	Size	Action
Week 2	Free Obsidian plugin, the 501-like live-preview fix	501 likes, 117 posts, open since 2022; Obsidian patches live preview monthly	Ship; kill under 200 installs in 14 days
Week 2 on	superpowers docs first, spec-kit second, plan-mode files third	22–30k, then about 1,660 active public projects a month	README line, MCP entry, "open the spec in frontmatter". Maintainer PRs are outbound; approval per PR
Now	Published research	Unpublished so far	The 157-source dataset
MVP-0 exit	Show HN	Zed's scored 43, Cursor's 9	One shot, after 6 of 10
MVP-1	Chrome extension, Product Hunt	1.5M browser readers	Backlinks, spike
MVP-2	Exported pages with a mark	No deploy count yet	Not a channel until measured

Users as evidence. The ten pilot strangers are the first interviews. Accept and revert events are logged in a local sidecar the user can read; no vendor usage database.

13 Risks

EXHIBIT 18

Ranked by cost if wrong

Risk	Response
Review state is a nice-to-have	Two-week test, three headlines; kill under 4 of 10 unprompted
Anthropic ships persistent review; plan-as-markdown exists since 2.1.70, diff review open at 262 reactions	Persistence across sessions, people and tools is what they do not build; re-check at week 12

Risk	Response
No team pays	20 conversations in weeks 1–2, on seats and gating
MVP-0 exit unmeasurable under the architecture	Written opt-in measurement decision before week 10
F5/F6 founder tests do not replicate with strangers	Include both in the ten-stranger pilot with a stated measure
Zed or Cursor add vault semantics	Depth in markdown and the state they do not keep
The 83% refusal is a class, not a bug	Patched corpus still refusing over 10 of 7,969 by day 24 falsifies it
Every AI call bills the operator	BYO key before any stranger
Client work crowds out product	Capped at 78 hours a month
GST	Reverse charge under CGST §24(iii) has no turnover floor

14 Ninety days and decisions




EXHIBIT 19

Fortnightly outcomes

Weeks	Work	Observable outcome
0	Rotate the two access tokens; decide the buyer; write the web byte path and the exit measure; attach the founder test notes for F5/F6	Written
1–2	Ten developers shown the review mock in both framings; plugin shipped; five people who bill for documents asked about approval; three landing pages; 20 team conversations; five people shown the S2 screenshot for five seconds and asked what the tint means	Kill signals at day 14
3	Go or no-go	Written decision
3–4	CI red run; NF-3, NF-1; content-hash review state prototyped on one spec-kit and one superpowers repo	Tint works on files frontmatter did not write
5–7	Review state store and rendering; BYO key; structured output; F5/F6	Stranger's key, stranger's repo, tinted spans
8–9	Revert, panel, answer in place, repo docs scan, engine wired behind every write	The fifteen-second demo
10–12	Polish; ten strangers	6 of 10 by the week-0 measure

EXHIBIT 20

Decisions

#	Decision	By
1	Review state as the product; attribution as a detail; F5/F6/F7 as MVP-0 features on founder evidence; the team tier as a hypothesis	The meeting
2	First buyer: AI-native developer, or the person liable when a document is wrong	Week 1
3	Share: hosted read-only viewer (a hosting obligation) or permalink plus export (no state; nothing for plain folders or private repos)	Week 2
4	Web byte path: browser-side fetch with the user's token, or through the server	Week 0
5	The two unrotated access tokens	Today

PART V

Research: interface, attention, retention, upsell, product family, format

15 Editor interfaces, benchmarked against shipped products

Scope of what was checked. The competitor teardown so far opened documentation, pricing pages and changelogs. A screen-level benchmark against shipped products through the Mobbin library was attempted on 2026-09-08 and **was not conducted**: every Mobbin call was refused by the session's injection-taint gate (the session had fetched pages carrying prompt-injection content; the gate blocks outbound connectors for the rest of that session, including read-only subagents). Nothing below is attributed to a Mobbin screen. The benchmark is scheduled as a week-0 task in a fresh session, with the query list in Exhibit 22.

EXHIBIT 21

What the incumbents' own documentation and trackers establish about interface placement

Question	Finding, from sources opened this session	Consequence for frontmatter
Is the AI affordance visible at rest?	No incumbent opens an AI panel expanded on load. Google Docs: <i>"By default, the Gemini bar automatically minimizes to clear your screen."</i> Notion: invoked by shortcut, space, or highlight, with an AI face on every page and no documented hide. Cursor: Cmd+K and Cmd+I. Zed and Obsidian Copilot: hidden until invoked. The incumbents' norm is a collapsed affordance present on load, not absence	The AI panel is one collapsed line at the bottom of S2; no model call on load
Do users want a global off switch?	Zed shipped <code>disable_ai</code> on 2025-07-23 after a 412-reaction issue; VS Code shipped <code>chat.disableAIFeatures</code> in 1.104 after a 30-reaction issue; Telegram hid its AI button in 6.7.4 after a 181-reaction, 154-comment issue	A hide-all-AI switch ships in S11 on day one and is named in the positioning
Where does review of agent output happen today?	Claude Code's VS Code extension opens plans as a full markdown document with inline comments (2.1.70, 2026-03-06); VS Code previews range-based feedback on markdown in an experimental agents window; the diff-review UI request is open at 262 reactions	The review panel (S4) is a persistent list, not a per-session diff; the tint is in the document, not in a side view

Question	Finding, from sources opened this session	Consequence for frontmatter
How do document tools scope AI context?	Five of six incumbents scope by file; Claude Code reduces a directory to "file listings, not contents"	S13 treats a folder as a pointer and lists what was read
What do editors ship in a first version?	Zed's launch page carried no AI; Obsidian's first public build had none	The free tier is the editor; AI is a panel, not the product

EXHIBIT 22

The screen benchmark to run in a fresh session (Mobbin, web, limit 6, batches of 2)

Maps to	Query
S0 launcher	"empty state for a new workspace with template cards to start from" · flow: "onboarding for a document or notes app that starts with opening or importing existing content"
S2 editor, AI placement	"Notion page with the AI writing assistant open, prompt input and suggestion chips" · "document editor with a collapsible page tree on the left, the document in the centre, and a properties or outline panel on the right" · "Craft document editor showing the sidebar and a page with blocks"
S2 right rail, S13 evidence rail	"Linear issue page with a right-hand properties panel and activity history"
S4 review panel	"code review screen showing a diff with accept and reject controls per change"
S11 settings	"settings page for AI features with an API key input field and a toggle to disable AI"
Pricing and upsell (§18)	section: "pricing section with a free tier and a team tier compared side by side" · "upgrade prompt shown when a user invites a teammate on a free plan"

Output required from that run: for each screen, the two or three distinct approaches found with a link to each; what every reference agrees on (convention) against where they diverge (a real choice); one borrow and one reject for frontmatter; and a ranked top five of changes to the prototype, each citing the reference that shows the better treatment.

16 Attention: what the evidence says about where to put things

No eye-tracking has been run on frontmatter's screens. Two primary sources apply, and one test in the plan closes the gap.

EXHIBIT 23

Evidence and the placement rules derived from it

Source	Finding	Rule for frontmatter
Tang et al., 2024, arXiv 2405.16081. Lab study, 28 developers validating Copilot-generated code, eye tracking plus IDE actions; one group told which code was LLM-generated, one not	"Without explicit information, developers often fail to identify the LLM origin of the code." Provenance awareness "led to improved performance, increased search efforts, more frequent Copilot usage, and higher cognitive workload"	The tint is justified by measurement: reviewers who know what the machine wrote validate better. The cost is workload, so the tint must be ignorable at rest (no border, no icon) and the detail must be on demand (hover, panel)
Nielsen Norman Group, F-shaped reading pattern	Fixations concentrate at the top and the left; "include the most important points in the first two paragraphs"	Every proposal card in S13 opens with its verdict and evidence on the first line. The review counter sits top-left of the panel. The file tree, where the unreviewed bars live, is on the left. The AI affordance is bottom and collapsed, where scanning does not reach until asked for
The five-second screenshot test, weeks 1–2 (§14)	Untested	Five people see the S2 screenshot for five seconds and say what the blue tint means. If they say "highlighted" or "error", the tint is wrong before any build

Predicted attention on S2, not measured. Top bar and left rail first (tabs, tree, the unreviewed bars); the document's first heading and opening lines second; the right rail last. The two things the product most needs seen, the tint and the review counter, are on that path. The thing it least wants to compete for attention, the AI panel, is off it. This is a prediction from the F-pattern, to be replaced by the five-second test.

17 Retention strategy

What is known. No comparable editor publishes retention. Accounts under \$50/month sit in the worst band (top-quartile annual gross retention 60–70%; 23% for AI-native products). Obsidian states it does not measure churn, and frontmatter's architecture is the same: no vendor usage database, so retention is observed only through opt-in, local, user-readable counts (§22, operating rules).

What the architecture allows. Files are the user's and exit is free, so retention cannot come from lock-in. It has to come from a reason to return.

EXHIBIT 24
Retention levers, each tied to a mechanism already in the plan

Lever	Mechanism	Why it fits
The counter	Every agent run leaves unreviewed spans; the launcher and the tree show the count. Opening frontmatter is how the count goes to zero	The return trigger is created by the user's own agent, daily, without any message from us
The scan	Repo docs scan finds stale sections and broken links after code changes	A second reason to return that does not depend on AI use at all
Answer in place and missing docs	Open questions and missing AGENTS.md/CHANGELOG are visible until closed	Small, finishable tasks
The plugin	The Obsidian live-preview fix stays installed in a vault the user already opens daily	Re-entry from a surface we do not own but they already use
First-session design	Open a folder, edit in thirty seconds, no account; the seven-concept budget	The first week decides retention; nothing is asked before value
No re-engagement email	Email is collected only on an obligation, never at first run	Retention is in-product or not at all; this is a constraint, stated
Measurement	Opt-in local counts: days-opened, spans reviewed, proposals accepted. The user can read the file	The MVP-0 exit test ("6 of 10 keep it") is observable only this way

18 Upsell and cross-sell

Every paid step is attached to a moment the product can see, and priced against what a comparable already charges for the same gate.

EXHIBIT 25
Upsell moments

Moment the product sees	Offer	Comparable gate	Tier
A second person opens the same repo	Shared review state: what each person has read; comments	GitHub Team, \$4, sells reviewers and code owners	Pro
A second device	Sync, provably safe	Obsidian Sync, \$4–5	Pro or Max
A repo with document gates wanted in CI	mdmax cert and review-state checks as a GitHub Action	Docs-as-code tooling; CI seats	Max
A user pastes a key, then asks for more than their free quota allows	Hosted AI, metered, hard-capped	Every AI product	Add-on, never default
The docs scan returns more proposals than one person will work through	A Zephyrus engagement: docs audit, AGENTS.md and spec-driven setup, review process design	The studio's fixed-scope engagements	Services

Cross-sell across the family. The editor is free and is the top of the funnel for everything else. The order is fixed: editor → team → CI → services. Each step uses the same sidecar and the same engine, so the upsell is more of the same thing across more people or more machines, not a different product. The one thing not cross-sold is anything outside the markdown core (§19).

Rules. No upgrade prompt before the moment exists. No nag: an offer appears once at the moment and lives afterwards in Settings. Downgrading loses features, never files. Unverified: whether small teams pay for shared review state at all; this is decided by the 20 conversations in weeks 1–2, not by the pricing page.

19 The product family, designed from the markdown core

The earlier ecosystem review (R22) asked which existing Zephyrus projects fuse with frontmatter and answered no for almost all of them, because they share a founder and nothing else. This section asks the other question: what family can be built around the one thing frontmatter owns, so that several things sell as one connected offering.

The coupling is a file format and an engine, not a brand. Every member reads and writes the same markdown, the same .frontmatter/review.jsonl sidecar, and goes through the same splice engine. That is what makes them one product to a buyer.




EXHIBIT 26

Family members, and what each sells

Member	What it is	Sold as	Status
Editor	The free product	Distribution	MVP-0
Team	Same sidecar across people and devices	\$4–5 per user per month	MVP-1/2
mdmax cert and gates	Renders the same certification across consumers, review-state and link checks in CI	Max, per team; the B2B line	MVP-2
MCP server	The write path for agents; carries prompts into the sidecar	Included; the integration that makes Team valuable	MVP-1
Chrome extension, Obsidian plugin	Read any .md with review state; the live-preview fix	Free; channels	MVP-1, week 2
Zephyrus services	Setting a team up: docs audit from the scan, AGENTS .md and spec-driven workflow, review process, migration from Notion	Fixed-scope engagements; the consultancy's productised entry point	Now, using the scan as the diagnostic
sgnk-md	The same codebase (202 of 228 files identical)	Merged, not sold separately	Parked
HQ, Markex, CareerOS, Brand OS, Advox, and the rest	Different buyers, different formats	Not coupled	Unchanged

The consultancy angle, stated plainly. The repo docs scan is a diagnostic that runs in the free editor and produces a list of things wrong with a team's documentation. Most teams will fix a few and ignore the rest. That list is the brief for a fixed-scope engagement, and the engagement leaves the team on frontmatter Team with CI gates. The product does the selling the studio cannot afford to do by hand. This is the one place where the AI-consultancy business and the product are the same funnel.

20 The markdown format: what MDMAX did, what it proves, and what "improvement" means

Verdict on formats, unchanged. The MDZ research settled it: no new format. Build a compiler and an editor for the markdown that exists. Every custom syntax degrades on some renderer; frontmatter adds nothing to a file that breaks it elsewhere.

MDMAX as built (from the line-level audit in docs/ENGINE .md): 3,614 lines across 13 files; 19 construct detectors over UTF-8 byte ranges with a skip mask; a degradation certificate that renders each block through 7 real engines against 15 (product, surface) targets and reports PASS, MUTATE or CORRUPT per cell; the byte-to-UTF-16 offset map; the splice writer. Verified on 8,513 pinned files. The artefact is a JSON sidecar, never written into the .md (proven).

EXHIBIT 27

What MDMAX proves, and what is still broken

Proven	Broken, measured
Byte-range splices leave every other byte identical, on 8,513 real files	The certificate CLI cannot start (ERR_MODULE_NOT_FOUND); one symbol of thirteen files reaches product code
The certificate never writes into the document	The consensus reference is wrong: on ~gone~ the engines split 3/2/2 and all seven cells report PASS; on task lists the four engines that do not implement GFM outvote the three that render correctly, so the correct ones are marked MUTATE
Callouts survive where fences do not: an unclosed fence swallows the document, a callout has no closer to lose (tested on four engines and GitHub)	certify drops the block after frontmatter when no blank line follows it; two fence grammars disagree inside one module
Tags and links can be written back in the file's own shape	OffsetMap .toU16 returns end-of-document at lengths that are multiples of 512, and reports ok: true

Proven	Broken, measured
8 of 15 target surfaces are uncertifiable and the code computes that share	6 of 19 construct detectors carry provenance : invented; six false-positive shapes reproduce

What "improving markdown" means here. Not syntax. Four things, all already in the engine's scope:

- 1. **Certified rendering.** A file passes when the engines that implement the spec defining each construct agree, not when a majority of arbitrary engines does. Replace the modal reference with a declared per-construct expectation named in `targets.ts`. This turns MDMAX into `mdmax cert`: a CI check that a team's docs render the same on GitHub, Obsidian and the consumers they name.
- 2. **Byte-safe writes, including frontmatter.** NF-1 through NF-4 are the writer's remaining refusals on real vaults; fixing them is the difference between 17% and near-100% of files the engine will touch.
- 3. **One sidecar beside the file.** Review state, comments and provenance live in `.frontmatter/`, plain text, anchored by byte range plus content hash. This is the only thing frontmatter adds, and it is not in the markdown.
- 4. **Non-English correctness.** The offset map is the seam; the D1 defect and the CJK tokeniser are the work.

Fix order for MDMAX (from the audit): make the CLI run (extensions and the `@/` alias, one day); run the certificate end to end once with real engines; land the four detector and offset tests that fail today; replace the consensus reference; then publish fidelity and refusal rates as one fraction, never fidelity alone.

PART VI

Reference

21 Rendering and markdown rules

Modes	Live, Edit, Split, Read. Nested constructs in list items render correctly (the 501-like bug, shipped first as the free plugin)
Rendered	Callouts > <code>[!kind]</code> ; Mermaid, rendered not extended; MADR and Nygard decision records; RFC 7322; Keep a Changelog 1.1.0; frontmatter as a properties panel
Not claimed	Runbooks and PRDs have no standard body; a structure is offered and labelled as such
Rule	Nothing is added to markdown that breaks it elsewhere; a touched file still renders on GitHub, in Obsidian, in a plain editor
Storage	Prose annotations as callouts (no closing marker to lose; an unclosed fence swallows the document). Tags and links written back in the shape the file uses. Review state and comments in the plain-text sidecar
Interop	Read <code>AGENTS.md</code> and <code>CLAUDE.md</code> , never replace them; MCP server for the user's agent; git as the only versioning; Obsidian conventions preserved

22 Operating rules

D2C / B2B	Free to build the audience; paid for teams. No enterprise feature until a customer refuses to pay without it: SSO at about 20 seats, SOC 2 above about 50 or for regulated buyers, a merchant of record for EU/US procurement
Communication	Email collected only when an obligation exists (payment, invite, recovery); never at first run. Breaking changes: two weeks' notice in product. Price changes: email before the next charge. All notices kept in an in-product list the user can check
Support	One address, a published response window; the free tier has no service commitment, stated on the page
Not measured	Session replay, keystroke telemetry, document content. Usage counts stay in a local sidecar
Mobile	About 15% of complaints; not in v1; later a reader and reviewer, not an editor
Accessibility, non-English	WCAG 2.2 AA; the <code>body-faint</code> token fails at 1.984:1 and is fixed in MVP-0; the tint is never colour alone. CJK: the byte map, a bigram search tokeniser, a written v1 scope decision
Security	No plugins, no code execution; the AI never gets ambient repo access; malformed repos refused, not repaired; 4 MB and 200,000-line ceilings per file

23 Where the plan stands

Strengths	The engine is correct where incumbents admit they are not; 88.6% of a shipping editor exists; six AI routes live; two founders funded by services
-----------	---

Weaknesses	Zero users; no CI, no BYO key, no structured output; 83% of vaults refused today; every channel borrowed; no evidence yet for the paid tier
Opportunities	Tens of thousands of active projects generating specs with nowhere to review them; review state is the one thing no incumbent persists; a 501-like bug fixable in a week; a hide-all-AI switch incumbents shipped only after backlash
Threats	Anthropic shipping persistent review; Zed or Cursor adding vault semantics in a sprint; Obsidian shipping first-party AI editing; attention running out before revenue