

frontmatter MVP 0: the build sheet

Approved. Decisions only, no argument. The evidenced version is the v3 plan.

Approved 16 September 2026. This is the version to print and build from. Decisions only, no argument. Numbers that carry a cost keep their source in brackets. Estimates say so.

Contents

01	What we are building
02	The twenty-two screens
03	Free and Pro
04	The blueprint
05	AI, credits, cost
06	Rendering
07	Storage
08	Offline, sync, live editing, desktop
09	Accounts, sharing, publishing
10	Money
11	Before the first stranger
12	Security
13	What we already have
14	The build
15	Risks, and where we stop
16	What we measure
17	Answer before phase 0 ends

§1 What we are building

An agent does better work when it is briefed properly. frontmatter writes the brief.

A markdown editor that turns an idea into a complete, consistent set of documents a coding agent builds from, and keeps them correct while it builds.

Sell it as: the best AI blueprints for your agentic development.

Who first: people who build with coding agents and start from an idea, not a repository. Then anyone who writes markdown and wants AI that respects the file.

The rule for the first five minutes. The editor is the landing page. No wall, no tour, no puzzle. One sign-in prompt, once, at the first action that needs it. The draft made before signing in comes along. Saving is never a button.

Where we win. Four platforms already generate plan documents free. None of them serves a checksummed set at an unauthenticated URL, none ships a map, and none keeps the documents editable in a real editor afterwards. Those are ours.

THE NINE ARTEFACTS AN AGENT NEEDS	WHO ELSE SHIPS IT
Brief and product document	Spec Kit, Kiro, ChatPRD, CodeGuide
Data and API contract	Kiro and CodeGuide, partly
Per-area specs	Spec Kit, Kiro
Verify commands inside each spec	Nobody
Decision records	Through a constitution, steering or rules file
Instructions file	Kiro, Cursor, Claude Code, Codex
Glossary	Nobody. Ours is MVP 1
A map of what connects to what	Nobody
Manifest with checksums	Nobody
Fetchable cold, no repository, no licence	Nobody
Editable afterwards in a real editor	Nobody. This is the product


§2 The twenty-two screens

Drawn to the shipped design system. Every value below was measured from the app on 16 September.

ELEMENT	VALUE
Header	52px, panel background, 14px side padding, 1px bottom border
The mark	26px square, 7px radius, accent fill, 13px at weight 800, wordmark beside it at 15.2px weight 650
Columns	264px left, flexible centre, 304px right, both side panes on the subtle ground
Tabs	38px tall, 13px at weight 450, right border, 2px accent underline when active
Toolbar	One row, 6 by 10px padding, 2px gaps, twelve buttons in a fixed order
Tool button	28px minimum width, 26px tall, 12px mono, transparent border
Mode control	Segmented: Edit, Live, Reading, Split, 12px at weight 500
Tree row	13px, 4 by 8px padding, 2px inset accent bar when active
Right pane	12px padding, uppercase 12px semibold section headings
Colours	<code>#fafafa</code> ground, <code>#18181b</code> ink and accent, <code>#0044cc</code> links, <code>#1a1a1a</code> dark ground
Shadows	None

Icons are Material Symbols, inline. Desktop 1440 by 900, phone 390 by 844. Source in `docs/mvp0/screens/`, one command to re-render.


S01. First run



S01 · The editor is the landing page.

Opens in under two seconds, empty, box focused. Every keystroke saved to this device. Export without an account. Four chips: blueprint, clean up a paste, write a spec, notes into a plan. A second row: open from GitHub, drop a .md file, a template. Any chip, AI edit, Share or Save to cloud opens S02.


S02. Sign in



S02 · Shown once, at the first action that needs an account.

Say what an account adds and what it does not change. Google first, GitHub second. The local draft moves in. "We never train on your documents", here and in the policy.


S03. The workspace



S03 · The shape everything else inherits.

Projects, folders, files, a badge for unread changes. Tabs coloured by project. Modes Edit, Live, Reading, Split. Twelve formatting buttons in the shipped order. Word count, bookmark, find, history, more. Right: outline always visible, then tags and bookmarks, backlinks, history (Pro), comments; Add file, Shortcuts, AI edit and the credits meter at the foot.


S04. The AI writing box



S04 · Only while the document is empty.

Appears on an empty document, disappears at the first character. The price in credits shown before anything is spent.


S05. Idea mode



S05 · Six decisions, each with a recommendation. "Not sure" takes it.

Four steps across the top. Left: the idea and the files that will be written. Right: six decisions, two or three options each, what each buys and costs, one marked recommendation.


S06. The blueprint, ready



S06 · A skill folder: index first, instructions, documents, specs, manifest, checksums.

Eleven files as a folder. `SKILL.md` is the index an agent loads first. `AGENTS.md` is emitted every time. The rail carries the file list, the consistency result in one sentence, the unlisted link, the kickoff prompt for the chosen agent, and the version. Publishing v2 writes a new path; a version never changes in place.


S07. AI edit



S07 · A proposal in the document. Accept or reject on the span.

AI never rewrites; it proposes. Old text struck, new text in the suggestion highlight the app already ships. The verb menu is the one already in the code: refine, expand, shorten, tone, translate, summarise, suggest links.


S08. Document review



S08 · Everything waiting, from people and from AI, in one list.

One list: collaborator suggestions and undecided AI edits, each with who, when, what, accept, reject. Accept all is one button. Comments sit beside it.


S09. Share



S09 · People need accounts. Publishing needs nothing.

Two halves. People, by email or Google account, can edit or can suggest. Publish, with a switch, the link, and the count used. Each cap states what Pro changes in one line, with no modal.


S10. A published page



S10 · Renders like the public page we already ship. No comments.

`noindex` by default with an owner switch. A dismissable "sign in with Google for a better view". A made-with line on free. Open in frontmatter, download .md. No comments, ever.


S11. Live collaboration



S11 · Presence, named cursors, the cap as a toast rather than a wall.

One collaborator on free. The second invite shows the upgrade.


S12. Plan and usage



S12 · Three meters, two plans, Team and Enterprise as coming.

Meters for AI edits with the blueprint credit beside it, cloud documents, published pages. Two plans with the caps from section 3. Team and Enterprise visible but unpriced.


S13. Offline



S13 · A banner, not a modal. The desktop prompt on a schedule.

Editing keeps working. Say where the text is and when it last synced. Surface a conflict; never merge silently. The desktop card appears after the third session, then every fourteen days until installed or dismissed twice.


S14. The desktop app



S14 · The same interface, signed, with a folder on this Mac beside the cloud projects.

One codebase, bundled. A local folder as a project. Files on disk written by the splice engine. Offline is complete here.


S15. Document history



S15 · Pro. Every save is a version; a version is a diff you can restore.

Versions with who and when, which were AI edits, a diff, restore, copy. Ninety days on Pro. Free sees the row and the pill.


S16. Custom blocks



S16 · A plain table, a chart block that reads it, Mermaid, a callout, maths.

Split view, source on the left and render on the right, plus a note on how each block looks elsewhere.


S17. The phone



S17 · Read, edit, outline, AI on a selection, share. Nothing else.

Installable. Chrome prompts on its own criteria; Safari needs an instruction sheet.


S18. Dark



S18 · The dark ground the app already ships.

Both themes ship. Nothing is light-only.


S19. Instruction files



S19 · The instructions file with a health panel and the agents that read it.

See and edit `AGENTS.md`, `CLAUDE.md` and rules files in the tree. Health panel: one file rather than two, size against Codex's 32 KiB cap, setup commands present, claims not verified recently. The agents that read it, listed. "Tidy this file" is one AI edit. The tool teaches the `@AGENTS.md` import; it does not diff two copies.


S20. The problems panel



S20 · Structural problems with line numbers, and one advisory writing note.

Broken wiki links, heading skips, missing alt text, table column mismatches, unknown front-matter keys. Structural checks run on the device and cost nothing. The writing note never blocks.


S21. A drawing on the idea



S21 · Attach a sketch, screenshot or repository before the questions.

The drawing is read once, described in words in the front-end spec, and never sent again, so the token cost is bounded and the privacy line stays honest.

S22. The map



S22 · What governs what, and which decision explains which spec.

Documents as nodes, what each governs as edges, decisions attached to the specs they explain. Counts that mean something: documents, links, orphans. Two files ship in the kit, `MAP.md` for a person and `graph.json` for an agent. Structure is read from the files, so it rebuilds on every save at no credit cost.

§3 Free and Pro

Free is free to try. Pro is do it properly. Every cap is server-side and changes without a release.

The editor, all free

Four modes. Twelve formatting buttons in the shipped order. Undo, redo, export to .md, HTML and PDF. Tree with projects, folders and an unread badge. Tabs coloured by project. Search across documents. Outline, tags, bookmarks, backlinks, properties. Shortcuts panel. Light and dark. Autosave on every keystroke. Problems panel.

Everything with a cap

FEATURE	WHERE	FREE	PRO
Documents in the cloud	Tree	10	Unlimited
Image uploads	In the document	5 MB a file, 100 MB an account	25 MB, 5 GB
Document history, diff, restore	S15	No	90 days
Trash	Tree	30 days	30 days
AI edits, eight verbs	S07	10 a month	100 a month
Careful mode, larger model	S07	2 credits	2 credits
Blueprints	S05, S06	1 a month	5 a month
Top-ups	S12	50 edits ₹99, 3 blueprints ₹149	Same
Live collaborators a document	S11	1	Unlimited
Comments and suggestions	S08	Yes	Yes
Published pages	S09, S10	5	Unlimited
Custom slug	S10	No	Yes
Made-with line on published pages	S10	Shown	Removed
GitHub repositories connected	S01, tree	1	Unlimited

FEATURE	WHERE	FREE	PRO
Pushes to a branch	S06	20 a month	Unlimited
Instruction-file editor and health panel	S19	Yes	Yes
Offline, PWA, desktop app	S13, S14, S17	Yes	Yes
The map, <code>MAP.md</code> and <code>graph.json</code>	S22	Yes	Yes

Integrations

Open any public GitHub markdown at `/gh/owner/repo/path.md`, no token needed. Import a Google Doc, capped at 10 MB by the Drive export. Import a Word .docx, converted in the browser. Serve every published page and kit as markdown at `page.md` and by `Accept: text/markdown`, plus `/llms.txt`. Open a folder on disk, including an Obsidian vault. Paste from ChatGPT or Claude and clean it up.

Not in MVP 0

Code review. Per-span "changed since you last read", which returns in MVP 1 as a diff over history. Book mode, slide mode, note insights. Multi-column layouts through raw HTML and CSS. An image puzzle anywhere. A student tier. A community beyond an opt-in index. Two-way GitHub sync, a glossary, an MCP server, a reverse blueprint from a repository and a signed Windows build, all MVP 1.

§4 The blueprint

booking-blueprint/	
SKILL.md	the index: what this is, when to load what
AGENTS.md	house rules every agent reads
00-BRIEF.md	the project in one page, the first user, the one metric
01-PRODUCT.md	what ships first, what is cut
02-DATA-AND-API.md	entities, relationships, operations
03-ARCHITECTURE.md	one short decision record per structural choice
04-SETUP.md	environment, local setup, the commands that gate a commit
specs/booking.md	files governed, verify commands, an exit condition
specs/payments.md	
MAP.md + graph.json	what governs what, which decision explains which spec
MANIFEST.json + SHA256SUMS	

Every `specs/*.md` carries verify commands and a numeric exit condition, so an agent can check its own work. The link is an unlisted address of at least 120 bits, `noindex`, `no-referrer`, revocable, never edited in place. A free kit unvisited for 30 days expires. `curl` needs no account.

§5 AI, credits, cost

Two credit types. An **edit credit** buys one AI edit, one generated document, one summary, one link suggestion. A **blueprint credit** buys one blueprint with its questions, writing and consistency pass.

Sonnet 5 is \$2 and \$10 per million tokens; Haiku 4.5 is \$1 and \$5. At ₹95.96 to the dollar. Token counts are assumed, so every line is an estimate.

TASK	TOKENS IN / OUT	SONNET	HAIKU
One edit	4,000 / 800	₹1.54	₹0.77
One document	2,000 / 1,500	₹1.82	₹0.91
One blueprint, eleven files plus a review pass	60,000 / 23,000	₹33.6	₹16.8
Blueprint, mixed: Haiku questions, Sonnet writing	20,000 / 3,000 then 40,000 / 20,000	₹30.2	

Edits run on Haiku by default, with a two-credit careful option on Sonnet. Blueprints run mixed. Prompt caching applies to the instructions, the question set and the templates, which are identical every time, so the marginal blueprint costs less than ₹30; measure it in week one.

At full use a free user costs about ₹37.9 a month and a Pro seat about ₹228, both estimates. Re-derive from measured usage after month one.

Ships with the credits: a daily spend breaker per user and one global, a per-task token ceiling, rate limits, and the six AI routes behind sign-in and metering.

§6 Rendering

A callout for prose, a fenced block for data, non-standard kinds prefixed `fm-`. No custom marker pairs: three hyphens already mean a thematic break and the front-matter fence, and an unclosed pair swallows the rest of the document.

The table stays a plain table; the block sits under it and reads it.

```
| Channel | Bookings |
|-----|-----:|
| Instagram | 312 |
| WhatsApp | 186 |
```

```
```fm-chart
kind: pie
table: above
```
```

Here the chart draws and the table folds. Anywhere else the table is still a table and the block is two lines of code. Kinds: pie, bar, line. "Convert to Mermaid" writes the portable version.

| BLOCK | HERE | GITHUB | OBSIDIAN |
|---|-----------------------|------------------------------|----------|
| Mermaid | Rendered, ships today | Rendered | Rendered |
| <code>fm-chart</code> over a table | Chart, table folded | Table then two lines of code | Same |
| Callout | Rendered, ships today | Rendered for five kinds | Rendered |
| Maths | Rendered, ships today | Rendered | Rendered |
| Tasks, tables, images, wiki links, properties | Ship today | Standard parts render | Rendered |

§7 Storage

R2 is the bucket, Firestore is the database, never the reverse. R2 has no query, no transactions and no bucket versioning. Firestore cannot hold a body past 1 MiB and charges per read.

| THING | WHERE | SHAPE |
|--|--|---|
| Document bytes, every version | R2 | <code>u/<uid>/d/<docId>/v/<n>--<sha256>.md</code> , never overwritten |
| Uploads | R2 | <code>u/<uid>/a/<sha256>.<ext></code> |
| Kit files, map, tarball | R2 | <code>k/<kitId>/v/<n>/...</code> plus checksums |
| Users, plan, credits ledger | Firestore, Mumbai | One ledger row per credit with task, model, tokens, cost |
| Index, projects, shares, publishes, comments | Firestore | Records only |
| Live session state | Durable Objects, one per open document | While the session lives |
| Security logs, 180 days | Mumbai | Legal floor |

Mumbai because the breach rules want logs held in India. Storage is not the cost; the model is. Move the R2 bucket to a Studio Zephyrus account before the first stranger's document lands in it.

§8 Offline, sync, live editing, desktop

Offline. This device first, cloud second. The service worker caches the shell. Saving compares versions: if the cloud moved, both are shown and the person chooses. Nothing merges silently. Safari deletes script-writable storage after seven days without interaction, so on Safari the promise holds only for the installed app. Chrome allows up to 60 percent of the disk.

Live editing. The document of record stays markdown bytes, one version per save. A live session is a shared document in a Durable Object that exists only while two people have the file open; every save writes a new version; when the last person leaves it is gone. It never owns the bytes, never decides a conflict between versions, never touches files on disk.

Desktop. Bundle the frontend, scope the file-system capability to folders the user picks, set a content policy, rename the identifier to something name-free, sign and notarise. macOS at \$99 a year. Linux needs no certificate but needs a Linux CI runner, because macOS cannot cross-compile. Windows waits for MVP 1. In-app version check, no auto-updater.

§9 Accounts, sharing, publishing

Google first, GitHub second, through the auth client already in the repository. Everyone who edits or comments needs an account, which is what makes suggestions attributable. Editing without an account exists only for the local draft and export. A published page is readable by anyone with the link, `noindex` by default, five free, custom slug on Pro, made-with line on free.

§10 Money

₹299 a month, ₹2,499 a year, India first. No dollar price in MVP 0.

A consumer price includes GST at 18 percent: $299 \div 1.18 = ₹253.39$. The gateway takes 2 percent plus GST, ₹7.06. **Net about ₹246 a month**, about ₹172 on the annual plan.

Card mandates register up to ₹15,000 without a fresh authorisation each time, and both prices clear it. UPI AutoPay needs a notice at least 24 hours before every debit, so the product emails first and offers a manual retry. An Indian card gets one attempt.

International payments are not a same-week switch: they need bank approval, video identity checks, and four published pages, terms, privacy, refund and cancellation, and shipping.

Carry both ₹299 and ₹399 into the first twenty conversations. Build at ₹299.

§11 Before the first stranger

| DUE BEFORE LAUNCH | WHAT |
|---|--|
| Privacy policy, terms, consent at sign-up | The 2011 rules are live now; the newer act's penalties start 13 May 2027 |
| Named grievance officer | The page and the address |
| Breach contact filed, six-hour runbook | Six hours from noticing, no size floor |
| Security logs, 180 days, in India | Minimal append-only log |
| Processor agreements | Cloudflare, Google, Anthropic, the payment gateway |
| Report form, 24-hour acknowledgement, takedown, preservation log, moderation page | Public pages by strangers make us an intermediary. About 74 founder-hours, and all of it ships before the first stranger publishes |
| Retention table, 18-plus line | One sentence now, awkward later |
| Four policy pages for international payments | Terms, privacy, refund and cancellation, shipping |
| EU sign-ups blocked | A representative costs €39 to €160 a month |
| GST position confirmed by a chartered accountant | Reverse charge on imported AI services |
| R2 bucket moved to the company account | Section 7 |

Never train on documents. Free kits expire. Delete on request. A published page is the only thing a stranger can read.

§12 Security

We render untrusted markdown, run a model over private documents, and write to GitHub. That is all three legs of the prompt-injection problem. Six controls ship in MVP 0.

1. **Untrusted content is never instruction.** Document text reaches the model inside a delimited data block with a standing rule that content inside it is data.
2. **No silent outbound fetch from a document.** Remote images in a shared document are proxied or click-to-load. The published page allows no third-party image or script origin.
3. **The AI proposes, never acts.** No AI-initiated share, push or publish, ever.
4. **A GitHub push is always explicit,** always to a branch, never a merge.
5. **The daily spend breaker** bounds an injection loop.
6. **No puzzle at the door.** Abuse checks run invisibly: app attestation on the AI routes, and an invisible challenge on anonymous publishing only if abuse appears.

§13 What we already have

| ALREADY BUILT | WHAT IT SAVES |
|---|---------------------------------|
| Complete payment integration: order verification, webhook signing, refunds, a browser test | Most of the payments item |
| The credit engine: plans, entitlements, billing tables, a three-tier AI cost model, kill switches | The ledger and server-side caps |
| A ten-primitive inline SVG renderer, one of them already a bar chart | Most of <code>fm-chart</code> |
| 19 browser-side PDF tools, no uploads | The browser export path |
| Two scoring engines, standard library only | The problems panel |
| A named-app auth client with a refresh guard | Sign-in |
| A numbered document convention proved on three sets, 61 files | The kit spine |
| A map of a whole system, rebuilt from files at no token cost | S22 |
| A public skills registry with 184 skills and a sync engine | The community route |

Genuinely new: the offline outbox, live editing, the GitHub App, the kit writer, the consistency check.

§14 The build

S is 1 to 2 engineering days, **M** 3 to 5, **L** 6 to 10. Every phase ends green, with a failing test first for anything that fixes a defect.

Phase 0, weeks 0 to 2, before code. Settle section 17. Move the R2 bucket to the company. File the breach contact, publish policy, terms and grievance pages, sign the processor agreements, ask the accountant, publish the four payment policy pages. Fix the repository default that still points at the sibling's vault. Publish the pace every Friday. Make twenty blueprints by hand for twenty people outside the studio and watch whether five run the kickoff.

| PHASE A, THE LAUNCH CORE | SIZE |
|--|------|
| Editor at the root with a local draft, no wall | S |
| Google and GitHub sign-in for strangers, a users record | M |
| Per-user storage: versions, index, projects, documents | L |
| Projects in the tree, tabs with project colours | M |
| Right pane: history and comments rows, shortcuts, AI edit, credits | M |
| AI writing box on an empty document | S |
| Export: wire the browser PDF path, delete the server route | S |
| Published page, noindex, sign-in card, made-with line | M |
| Report form, 24-hour acknowledgement, takedown, moderation page | M |
| Markdown twins and <code>/llms.txt</code> | S |
| Problems panel | M |
| 27 to 46 days | |
| PHASE B, AI AND CREDITS | SIZE |
| Suggestion mode: proposals in place, accept and reject per span | L |
| Credits ledger, two types, caps, breaker | M |
| Routes behind sign-in and metering, model routing, prompt caching | S |
| Injection controls: delimited data block, image proxy, page policy | S |
| 11 to 19 days | |

| PHASE C, THE BLUEPRINT | SIZE |
|--|------|
| Six decisions from a prompt, in the card shape | L |
| Kit writer: eleven files, one review pass, consistency check | L |
| Skill folder, index file, instructions file, verify arrays | S |
| The map: <code>MAP.md</code> and <code>graph.json</code> , rebuilt on save | M |
| Storage, link, manifest, tarball, checksums, revoke, expiry | M |
| Kickoff prompt, agent picker, copy | S |
| A kit opens as a project; edit then publish v2 | S |
| Attach a drawing, document or repository to an idea | M |
| GitHub App: connect, open a repository, push to a branch, caps | L |
| 30 to 51 days | |

| PHASE D, SHARING AND COLLABORATION | SIZE |
|--|------|
| Share dialog, roles, invites by account | M |
| Comments on spans | M |
| Suggestions from collaborators into the review panel | M |
| Live editing: shared session per document, presence, the cap | L |
| 15 to 25 days | |

| PHASE E, OFFLINE, PWA, DESKTOP | SIZE |
|--|------|
| Local-first outbox, compare-and-swap save, conflict prompt | L |
| PWA: cached shell and documents, icons, install prompt, Safari sheet | S |
| Desktop nudge on its schedule | S |
| Bundle, file-system scope, content policy, identifier, macOS signing and notarisation, version check | L |
| Linux targets on a CI runner | S |
| 15 to 26 days | |

| PHASE F, PRO AND THE REST | SIZE |
|---|------|
| Subscriptions, webhooks, GST invoices, one-attempt retry | S |
| Plan and usage page, upgrade, top-ups | M |
| Document history, 90 days, diff, restore | M |
| <code>fm-chart</code> over a table, convert to Mermaid, elsewhere preview | M |
| Instruction-file editor with the health panel | M |
| Imports: Google Docs and Word | M |
| Retention table, 18-plus line, EU block | S |
| 17 to 29 days | |

The sums. Everything: **115 to 196 engineering days**. The launch core, what a stranger meets on day one, is phases A, B and C plus payments and the plan page: **72 to 123 days**.

| AT THIS PACE | LAUNCH CORE | EVERYTHING |
|-------------------------------------|-----------------|-----------------|
| 1.21 days a week, the measured rate | 60 to 102 weeks | 95 to 162 weeks |
| 3 days a week | 24 to 41 weeks | 38 to 65 weeks |
| 5 days a week | 14 to 25 weeks | 23 to 39 weeks |

If the date slips, cut from the bottom: live editing, then the desktop build, then GitHub, into releases after launch. The PWA covers the download prompt.

§15 Risks, and where we stop

| RISK | KILL LINE |
|--|--|
| Nobody keeps a blueprint alive after the first build | Fewer than two of ten pilot users edit a kit after its first build: keep the editor, stop the generator |
| Free AI costs more than assumed | Model cost per active free user over ₹60 in a month: halve the free caps that week |
| Pro loses money at full use | Measured cost per Pro seat over ₹200: ₹399 for new seats |
| Three unfamiliar technologies at once | If the live-editing spike runs past ten days, buy it instead |
| Injection through a shared document | One confirmed leak: sharing goes invite-only until fixed |
| A breach in the first month | Six hours to the authority; the runbook is written before launch |
| A platform absorbs the kit | If a major platform ships a public shareable project link, the pitch narrows to the editor and the desktop app |
| Pace | Launch core not done in 30 weeks: cut blueprint automation, launch the editor with hand-made kits |

§16 What we measure

The funnel: visits, documents created, sign-ups, blueprints started, finished, kickoff copied, kit fetched by an agent, **kit edited again within seven days**, pages published, second collaborator invited, Pro.

The one number: blueprints edited after their first build, per week. It tests the only claim that separates us from every generator.

Also: the acceptance rate of AI suggestions, accepted against proposed.

The pilot passes when three of ten strangers name the problem unprompted, two of ten are still editing at day 30, one call is booked, and five of ten run the kickoff.

Published every Friday: engineering days, the running rate, the re-forecast date, model cost per active user, suggestion acceptance, and conversion once it exists.

§17 Answer before phase 0 ends

1. **The pace, and who builds.** At the measured 1.21 days a week the launch core is a year out. The largest decision in this file.
2. **₹299 or ₹399.**
3. **The name.** No trademark search has succeeded: the registers need a login and a puzzle. Free if we coin: `getfrontmatter.com`, `usefrontmatter.com`, `mdmax.in`, `mdmax.ai`, and the `mdmax` package name.
4. **Windows:** the PWA now, a certificate in MVP 1?
5. **EU sign-ups** blocked for MVP 0?
6. **The templates source.** No "box project" repository exists on this machine. Default is the eleven-file kit in section 4.
7. **The free caps** as written?
8. **Who makes the twenty hand-made blueprints, and by when.**
9. **Which of the 184 skills may be public.**
10. **The map's role:** a view plus two files in the kit, and does the kickoff prompt tell the agent to read it first?
11. **Project colour:** chosen or assigned?
12. **The kit shape for larger projects:** one folder, or one per area?